

Training data for AI, *taylor-made.*

Artificial intelligence needs data to learn. We source it, prepare it, and deliver it ready to use.

So your team builds better models, faster.

— THE CORE IDEA

AI is the chef. Data is the ingredients.

Without good ingredients, even the best chef can't make a good dish. Today almost everyone is competing for the best chefs, while almost no one is taking care of the supply. That's where Caudals comes in: we bring the raw material AI needs to learn.

— CONTENTS

Five chapters.

01	A problem in two parts	MARKET
02	What we do and how it works	PRODUCT
03	Who buys and who sells	CUSTOMERS
04	Market position and model	BUSINESS
05	Traction, roadmap, and the ask	CLOSING

A problem in two parts.

AI teams can't find the data they need. The companies that have it can't sell it.

— THE BUYER SIDE

AI learns by looking at examples:
without good examples, there is
no good AI. And good examples
are not easy to come by.

The question today isn't which model we train.
The question is what we train it on.

T00 GENERIC

Public data doesn't fit the sector, geography, or format the buyer actually needs.

T00 AMBIGUOUS

Usage rights are unclear. Legal won't approve training on it.

T00 SLOW

Internal teams and labelers can't source, license, and deliver end to end.

— WHY PUBLIC DATA ISN'T ENOUGH

You can't build a real edge on data your competitors have downloaded too.

3_x

more model performance, on average, with sector-specific data versus public data.

72%

of AI teams say their biggest obstacle is finding proprietary data with clean permissions.

8_{wks}

average time it takes a buyer to find usable data before training can begin.

Indicative figures based on interviews with buyers across logistics, retail, healthcare, and finance.

— THE PROVIDER SIDE

Thousands of companies have highly valuable data sitting in storage, and don't know how to sell it.

A	They have what AI wants Transactions, routes, sensor readings, operational logs, inspection reports, support transcripts, business archives.	WHAT EXISTS
B	Selling it means standing up another company Sales team, legal review, privacy, technical formats, documentation, quality control, buyer network, and contracts.	WHAT IT TAKES
C	So the data stays in the drawer Opportunity cost piles up. Some companies try once and give up. Most never try at all.	WHAT HAPPENS

What we do.

Caudals sits in the middle, between the companies that have the data and the teams that need it.


— THREE CAPABILITIES

Provider sourcing, data preparation, and commercialization.

01	Provider sourcing We identify companies whose proprietary data is useful for training a model. We validate rights, negotiate access, and scope the opportunity.	SOURCING
02	Data preparation We turn raw files into a training-ready dataset. Ingestion, cleaning, anonymization, enrichment, labeling, quality control, and packaging.	ENGINEERING
03	Commercialization We verify, list, license, and deliver. We negotiate with the buyer. We share revenue with the provider on every sale.	TRANSACTION

THE DATA PREPARATION PIPELINE

Seven steps. From raw files to AI-ready data.



Every dataset we deliver runs through this. **No one else in the market operates the full chain** from source to delivery.

TWO SIDES, ONE PLATFORM

Providers on one side. Buyers on the other. We operate the middle.

PROVIDERS · COMPANIES WITH DATA

What they have

- Transactions, invoices, receipts
- Sensors and IoT devices
- Logistics and routes
- Retail and commerce behavior
- Operations and inspection logs

REVENUE · SALES SHARE

BUYERS · AI TEAMS

What they need

- Data to train and fine-tune
- Evaluation sets
- Edge and adversarial cases
- Recurring data feeds
- Domain-specific corpora

RECEIVE · AI-READY FORMATS



Who we serve.

Buyers whose next model is stuck on data. Providers whose data sits idle without a path to market.

— WHO BUYS

AI teams whose next model is blocked by the absence of data.

COMMON PROFILES

- AI startups training in a specific sector
- Machine learning teams at large enterprises
- Product teams fine-tuning language models
- Computer vision teams that need image and video
- Teams in logistics, retail, healthcare, agriculture, and finance

WHAT THEY SAY

"The model is designed. We've spent four months hunting for the right data to train and evaluate it."

Head of AI, Series B fintech

— WHO SELLS

Companies that don't see themselves as data vendors.

TYPICAL SECTORS

LOGISTICS

Routes, ETAs, fleet telemetry.

INDUSTRY AND IOT

Sensors and maintenance.

HEALTHCARE

De-identified records and imaging.

RETAIL AND COMMERCE

Transactions, baskets, returns.

AGRICULTURE

Harvest, satellite, environmental.

SUPPORT AND DOCUMENTS

Tickets, transcripts, KBs.

WHAT THEY SAY

"We have fifteen years of operational data sitting in storage. And yet we have no idea how to sell it without standing up an entire team to do so."

Chief Data Officer, national logistics operator

Where we fit.

Four adjacent categories each solve part of the problem. We combine them and add what none of them does.

Four categories. Each solves only one challenge.

CATEGORY	DEMAND	SOURCING	LICENSING	OPERATIONS	REV. SHARE
Data marketplaces AWS Data Exchange · Snowflake · Databricks · AfterQuery	✓	—	—	—	partial
Labeling and annotation Scale AI · Appen · SageMaker Ground Truth · Mercor	—	—	—	✓	—
Data and AI consultancies Custom pipelines and projects	—	✓	partial	✓	—
Monetization platforms Packaging for providers	partial	—	partial	partial	✓
Caudals Marketplace and managed operations	✓	✓	✓	✓	✓

SERVICE TIERS

Three ways to work together.

TIER 01

Feasibility study

\$1,000

5 TO 10 BUSINESS DAYS

Feasibility study, risk map, target schema, and a go/no-go recommendation.

CREDITS AGAINST THE PILOT

TIER 02

Dataset pilot

Custom pricing

2 TO 4 WEEKS, 5,000 TO 50,000 RECORDS

Schema, cleaning, anonymization, light labeling, QA, and a sample.

CREDITS 30–50% AGAINST THE PROJECT

TIER 03

Full dataset

Custom pricing

6 TO 20+ WEEKS, PRODUCTION QUALITY

Multi-source sourcing, full pipeline, catalog-ready packaging, and final delivery.

OPTIONAL CATALOG LISTING

Fees today. Compounding effect over time.

REVENUE

Services

Fixed and milestone-based fees on feasibility studies, pilots, and full projects.

Transactions

Revenue share with the provider on private offers and items in our catalog.

The catalog compounds

Every completed project becomes a resellable asset. A dataset delivered once can be sold many times over.

LIVE TODAY

Public site, contact, blog, and waitlist

Admin panel with leads, builds, and operations

Pipeline compliant with data protection regulations

Available in Spanish and English

• Onboarding buyers

• Onboarding providers

— FOUNGING TEAM

Two complementary profiles. Business and technology.



CO-FOUNDER · CEO

Mario Medrano Paredes

Leads strategy and the data preparation pipeline. Ingestion, cleaning, anonymization, labeling, quality control, and delivery. Owns datasets and privacy.

Telecommunications Engineer. Prior background in data analytics, machine learning, and AI models.



CO-FOUNDER · CTO

Alonso Sandoval Martínez

Leads two-sided acquisition, data provider partnerships, buyer relationships, and operations. Owns technology and commercial.

Telecommunications Engineer. Prior background in banking application development, cloud platforms, and B2B sales.

Let's talk.

AI is moving faster than the data that feeds it. The gap between those who hold the data and those who need it is a singular opportunity. If this side of the market interests you, we would love to explore ways to work together.

TODAY TO 3 MONTHS

Operationalize the middle

Customer acquisition, provider onboarding, tracking dashboard, pilot projects

3 TO 6 MONTHS

Public catalog

Catalog items, private offers, recurring data feeds, self-service

6+ MONTHS

AI data marketplace

The leading platform for ready-to-train AI datasets across every sector

CONTACT

hello@caudals.com · caudals.com