

How to pivot Caudals and land your *first* paying customer.

A narrow product definition, a demonstrable artifact, a repeatable way to earn trust without a reputation, and a 90-day plan built for two people with a VPS.

THE THESIS IN ONE SENTENCE

Stop selling *datasets*, which nobody knows how to buy from you, and start selling **proof that an AI system is wrong** — then sell the data work that fixes it.

PREPARED FOR

Mario & co-founder

BASED IN

Valladolid, Spain

HORIZON

90 days to first revenue

DATE

September 2026

What is in this document

Read sections 01–05 first; they contain the decision. Sections 06–10 are the go-to-market. Sections 11–13 are the build. Everything after 14 is reference material you will use week by week.

01	The decision, on one page	What to do Monday, compressed.
02	Diagnosis: why you have been stuck	Three structural errors, not a motivation problem.
03	What changed in the market in 2026	The window that makes this pivot work now.
04	The pivot	Same pipeline, different entry point.
05	The product, precisely	Four SKUs, what is in the box, what you never promise.
06	Alternatives considered	Five options scored honestly, with the losers explained.
07	Who you sell to	Two sectors, one trigger, and how to build the list.
08	The Reality Check play	Earning trust with evidence instead of a pitch.
09	Landing the first customer in 30 days	Day-by-day, with the founding-customer offer.
10	Pricing and the revenue ladder	€0 → €2,400 → €890/mo → €18,000.
11	Technical implementation	Schema, runner, scoring, and the public demo.
12	Keep, freeze, kill	What to do with everything already built.
13	The 90-day roadmap	Weekly milestones with explicit gates.
14	Metrics and unit economics	The eight numbers that matter.
15	Legal, ethical and compliance guardrails	Non-negotiable rules for testing other people's systems.
16	The marketing engine	What to publish, forever, without inventing topics.
17	Risks and failure modes	How this dies, and the antidotes.
18	Operating cadence for two people	Role split, weekly rhythm, decision rules.
A	Outreach templates	Spanish and English, ready to send.
B	Eval case schema and scoring rubric	The technical contract of the product.
C	The report deliverable	Exact table of contents of what you hand over.
D	Target-list construction	Queries and sources to find 40 companies in a day.
E	Sources	Everything factual in this document, referenced.

On one page

If you read nothing else, execute this. Every other section exists to justify, sharpen, or support it.

The move

Caudals stops presenting itself as a data marketplace and becomes an **AI reliability company** for the Spanish mid-market: you test whether a company's deployed AI assistant actually answers correctly, you deliver a scored, evidence-backed report, and you sell the ongoing test suite and the data work that closes the gaps.

You do not abandon the dataset business. You *invert its order*. Today you ask companies to buy a dataset before they can see anything. Instead, you hand them a measurement of their own failure — free, in 48 hours, unrequested — and the dataset work becomes the obvious next purchase they ask you for.

OLD SEQUENCE

Pitch marketplace → explain what a dataset is → propose €12k pilot → prospect stalls → nothing exists to show → months pass.

NEW SEQUENCE

Test their live assistant → send three real wrong answers → free full report → €2,400 test suite → €890/month retainer → €18k data repair build.

The first product

Caudals Evalset: a 150–300 case golden test set built from the customer's own documentation and their real customer questions, run against their assistant, scored on accuracy, grounding, false confidence and safe refusal, delivered as a live dashboard, a signed PDF report, and a portable `JSONL` file they own. Price: **€2,400**. Delivery: two weeks. Recurring: **€890/month** to re-run it monthly and track regressions.

The first customers

Spanish companies of 50–500 employees that **publicly launched an AI assistant in the last 18 months** in two sectors: (1) insurance, health insurers and brokers; (2) industrial after-sales and technical support, which is dense around Valladolid. Public launch is the trigger: it proves budget was spent, someone owns it internally, and the system is exposed to customers where errors are expensive.

The unlock for your credibility problem

You cannot borrow reputation at 23 with no logos. You can *manufacture* it. A report containing seven verifiable wrong answers from their own production system is not an opinion, and it does not care how old you are. Evidence outranks credentials. This is the single most important idea in this document.

DO THIS ON MONDAY

- Choose **one** sector. Insurance, unless you already have a warm contact elsewhere.
- Build a list of 40 Spanish companies in that sector with a live, publicly reachable AI assistant (Appendix D shows how, in about six hours).
- Manually run 40 questions against three of them. Spreadsheet only. No code.
- Write one 6-page report by hand. Make it beautiful. That artifact is your entire company for the next 30 days.
- Do not write a line of platform code until you have delivered three of these reports.

€2,400	30	3	0
Target first invoice, small enough for one director to sign without procurement	Days from starting the list to the first signature	Founding customers at a discount in exchange for named case studies	Lines of new platform code required before the first sale

Why you have been stuck

You have not been failing at execution. You have been executing correctly against a business model that cannot produce a first customer at your size. Three structural errors explain every symptom you described.

Error one — you are selling an input, not an outcome

Nobody has ever woken up wanting a dataset. They want a model that behaves, a claim that gets classified correctly, a support cost that goes down, an audit they pass. A dataset is an ingredient. Selling ingredients requires the buyer to already have a recipe, a kitchen, and a chef — that is, a mature ML team that knows exactly what data gap it has. In Spain, outside a handful of large banks, telcos and IBEX-scale companies, that buyer barely exists, and those that do exist buy from Scale, Surge, Appen, Mercor or their own internal teams.

The Spanish small and mid-market company you can actually reach has the opposite problem: it has *already deployed* something with AI in it, and it does not know whether it works. That is a named, felt, budgeted pain. It is also, mechanically, a data problem — but you must never introduce it as one.

Error two — your smallest sellable unit is far too large

Your current ladder starts at a \$1,000 feasibility study and jumps to a \$12,000–\$35,000 pilot. For a company that has never heard of you, run by two people with no track record, both are enormous asks. A €1,000 study is *simultaneously* too expensive to be an impulse and too small to justify a procurement process — the worst possible position on the price curve.

Worse, both products are promises. The buyer pays first and sees the artifact later. With zero reputation you are asking for a leap of faith you have not earned and cannot yet earn.

Error three — you have no evidence surface

Everything a prospect can currently see about Caudals is *asserted*: a landing page, a blog, a description of a pipeline. There is nothing they can touch, run, or verify. Your instinct in the brief — "we need something tangible, something they can visualise and test" — is exactly right, and it is the highest-leverage fix available to you.

Mapping your stated problems to their real cause

WHAT YOU EXPERIENCE	REAL CAUSE	WHAT THE PIVOT DOES ABOUT IT
Large companies request enormous, specialised datasets (e.g. labelled medical data) that we cannot source or process.	You positioned as a supply-side vendor, so buyers hand you their hardest sourcing problems.	You stop sourcing entirely. The raw material becomes the <i>customer's own</i> documents and question logs, which they hand you on day one.
Big, established companies do not take us seriously — two young founders, no reputation.	You are asking for trust <i>before</i> demonstrating competence. Credentials are the only substitute, and you have none.	The first contact <i>is</i> a demonstration of competence: specific, checkable errors from their own system. Age becomes irrelevant.
We do not know where to get data or how to complete every processing step.	Unbounded scope. "All data, all modalities, all sectors" has no learnable craft.	One modality (text and documents), one artifact (a scored test suite), one craft you can master in six weeks.
Small Spanish companies do not understand what a dataset is or why they need one.	You are asking a non-technical buyer to learn your category before it can buy.	You never use the word. You say: "your assistant answered 34 of these 100 real customer questions incorrectly. Here they are."
Only a VPS for infrastructure; large-volume processing is hard.	The old product's cost scales with data volume, which you cannot finance.	A 300-case test suite is a few megabytes. Your existing VPS is already over-provisioned for this product by an order of magnitude.
Nothing to demo; prospects go cold.	No artifact exists before a large contract is signed.	The artifact exists <i>before</i> the first conversation, and a public self-serve demo exists on your own domain.

THE HONEST PART

Your build is not wasted. You have a Next.js platform, Supabase, object storage, an admin console, a leads CRM, a blog and newsletter pipeline, and deployment discipline. That is roughly nine months of infrastructure most people at your stage do not have. The problem was never the machine — it was that the machine was pointed at a market requiring capital and reputation you will not have for two years. Section 12 says exactly which parts to keep, freeze and shut down.

What changed in 2026

This pivot is not a generic "do consulting instead" suggestion. Four specific shifts in the last eighteen months created a gap that a two-person team in Valladolid can credibly fill right now, and that gap will close within roughly two years.

1. The centre of gravity moved from labelling to evaluation

Through 2023–2024 the money in AI data was in annotation volume. In 2025–2026 it moved decisively toward **environments and evaluation** — data that measures and trains behaviour rather than data that merely describes the world. Reports in late 2025 described Anthropic weighing over \$1 billion in spend on reinforcement-learning environments; in February 2026 Mercor acquired Sepal AI specifically to deepen environment capability, and Scale, Surge, Turing and Centific now all sell evaluation and environments alongside classic human data.

That is the frontier-lab tier and it is closed to you. But the same logic is diffusing downward: every company that deploys an assistant eventually discovers it needs a test suite. At the mid-market almost nobody is serving that need, because it is too small for the giants and too unglamorous for AI startups chasing product-led growth.

2. Deployment massively outran reliability

The consistent finding across 2026 industry surveys is a large gap between pilots and production. Composio's agent report found roughly two thirds of organisations reporting measurable gains from pilots but only around one in ten scaling to production; a March 2026 survey put enterprise agent piloting near 78% with production deployment below 15%; Dynatrace's 2026 Pulse found roughly half of programmes still stuck in proof-of-concept, citing reliability and observability as the gating factors.

Read that as a market signal rather than a statistic: there is a very large population of companies that *have already spent money* on AI and cannot demonstrate it works. They are your customers. They are not waiting to be educated about AI — they are waiting for someone to tell them the truth about what they bought.

3. Spain specifically is adopting fast and measuring badly

Spanish INE data puts AI use at 21.1% of companies with 10 or more employees, up from 12.4% the previous year — one of the sharper growth rates in the EU. A YouGov study for IONOS found 35% of Spanish companies planning AI budget in 2026, up from 22% in 2025. Meanwhile the European Commission's Digital Decade reporting placed Spanish SME adoption at 10.3% against an EU average of 13.5%.

The combination that matters: **fast adoption, thin internal capability**. Companies are buying assistants from integrators and agencies, and neither buyer nor vendor has an objective way to say whether the result is good. There is no independent referee in this market in Spain. That is the position you should take.

21.1%	35%	<15%	10.3%
Spanish companies (10+ employees) using AI, up from 12.4% previous year	Spanish companies planning AI budget in 2026, from 22% in 2025	Enterprise agent pilots reaching production — 2026	Spanish SME AI adoption vs 13.5% EU average — EC Digital Decade

4. Regulation is now a scheduled, dated obligation — later than you may think

Be precise here, because getting it wrong destroys credibility in a sales conversation. The EU's Digital Omnibus — Regulation (EU) 2026/1744, published 24 July 2026 and in force from 27 July 2026 — **deferred the high-risk obligations**. Stand-alone Annex III high-risk systems now face **2 December 2027**; AI embedded in regulated products under Annex I moves to **2 August 2028**. The August 2026 date that circulated widely for the last year is no longer the operative deadline.

In Spain, the Council of Ministers approved the draft Organic Law on the good use and governance of AI on 26 May 2026; it entered Congress and was still in parliamentary process through the summer. It designates AESIA (headquartered in A Coruña) as the supervisory agency, establishes regulatory sandboxes, and sets a national sanctions regime with fines reaching €35m or 7% of global turnover for prohibited practices.

HOW TO USE REGULATION IN SALES — AND HOW NOT TO

Do not lead with compliance fear. The deadline moved, buyers know it moved, and fear-selling a delayed regulation makes you sound like every other consultancy sending the same email. It also positions you as a legal advisor, which you are not.

Do use it as the second reason, framed as timing: the Article 10 data-governance and Article 15 accuracy and robustness expectations require documented evidence about training, validation and test data — relevance, representativeness, error rates, provenance, annotation, gaps. A company that starts building a test suite in 2026 arrives at December 2027 with eighteen months of run history. A company that starts in mid-2027 arrives with nothing. Sell the head start, not the panic.

What this means for your window

Independent AI evaluation for the European mid-market is currently under-served, unbranded, and unregulated as a profession. Within roughly two years, consultancies and the Big Four will productise it, and evaluation platforms (Braintrust, LangSmith, Arize, Opik, Langfuse and others — a category that has already seen significant consolidation) will push down-market with self-serve offerings. Your advantage is not technology. It is that you will do the unglamorous human work of writing correct answers with a domain expert, in Spanish, for a company in Valladolid, for €2,400 — and nobody larger wants that job.

Same pipeline, different entry point

You asked not to deviate too far from the original idea. You will not. Every stage of the Caudals pipeline survives this pivot intact; what changes is where the raw material comes from, what the output is called, and who pays for it first.

The pipeline, remapped

ORIGINAL CAUDALS STAGE	BECOMES	WHY IT IS NOW EASY
Sourcing & supplier acquisition	Customer's own documents, FAQs, manuals, policies, ticket logs	The customer hands it to you in week one. No sourcing, no negotiation, no acquisition cost.
Rights & licensing review	A one-page data-processing annex in the contract	They own the data. Scope is a DPA and an NDA, not a licence negotiation.
Profiling	Coverage analysis: which topics the knowledge base does and does not answer	This becomes a headline finding in the report, not internal work.
Cleaning & normalisation	Question canonicalisation, deduplication, ambiguity resolution	Hundreds of records, not millions. Runs on a laptop.
PII handling	Redaction of real customer questions before they enter the suite	Small volume, mechanical, and a visible trust signal in the deliverable.
Labelling & curation	Authoring the reference answer, the accepted variants, the citation, the difficulty tier	This is the actual product. It is where your craft lives and where you get better than competitors.
QA & scorecard	The scored run: accuracy, grounding, false confidence, refusal correctness	Turns into the number the customer stares at. The most tangible thing you will ever ship.
Packaging & delivery	Dashboard + PDF report + <code>JSONL</code> golden set the customer keeps	The <i>JSONL</i> is a dataset. You end up selling exactly what you always wanted to sell.
Provenance & audit	Every case cites its source document, version and author; every run is immutable	Directly reusable as AI Act Article 10/15 evidence in 2027.

THE SENTENCE THAT HOLDS THE WHOLE STRATEGY TOGETHER

"We don't sell you data. We show you exactly where your AI is wrong — and then we build the data that fixes it."

Why measurement is a better wedge than data

1. **It is self-evidently true or false.** A wrong answer is a wrong answer. There is no need to persuade anyone of the value of a dataset in the abstract.
2. **It creates its own demand for the next product.** Every failing case is a data gap with a price tag. You do not need to invent a reason to propose the €18k build — the report proposes it for you, in the customer's own words.
3. **It is repeatable revenue by nature.** A test suite is worthless run once. Models change, prompts change, documents change. Monthly re-runs are not an upsell tactic; they are the honest shape of the product.
4. **It is cheap to produce and expensive to ignore.** Your cost is human hours; their cost of being wrong in front of a customer is reputational and sometimes regulatory.
5. **It is the only position from which a nobody can approach a somebody.** You are not asking for their time. You are giving them information about themselves that they do not have.

What you are explicitly giving up

Say this out loud to each other, so it is a decision and not a drift:

- The two-sided marketplace. Deferred indefinitely, not "next quarter". A marketplace needs liquidity on both sides; you cannot fund that.
- Supplier revenue-share and Stripe Connect. Frozen.
- Image, video and sensor modalities. Text and documents only for at least six months.
- Large-volume data engineering. If a deal requires processing more than a few gigabytes, it is out of scope in 2026.
- Selling to frontier labs and IBEX-35 enterprises. Not forever — just not as the path to customer number one.

The brand, the domain, the platform, the content engine and the eventual ambition all survive. You are narrowing the beam, not changing the target.

The product, precisely

Four SKUs forming a ladder. Each one exists to sell the next. Prices are in euros, calibrated to Spanish mid-market signing authority, and deliberately low for the first six months — you are buying case studies, not maximising margin.

SKU 01 · LEAD GENERATOR

Reality Check

€0 — 48 hours

What it is: 40 questions run against a company's publicly reachable AI assistant, scored, with every failure documented as a reproducible transcript.

Deliverable: a 6-page PDF. Score out of 100, five failure categories, seven annotated transcripts, one page on what it would take to fix.

Purpose: this is not a product, it is your business card. It replaces the cold email entirely.

SKU 02 · FIRST INVOICE

Evalset

€2,400 — 2 weeks

What it is: 150–300 case golden test set built from their documentation, their real customer questions, and one 90-minute session with their domain expert. Run against their assistant. Scored.

Deliverable: live dashboard, 20-page report, and the `JSONL` golden set they own outright and can take anywhere.

Purpose: the first invoice. Small enough for one director to approve.

SKU 03 · THE REAL BUSINESS

Reliability Retainer

€890/month — 12-month term

What it is: monthly re-run of the suite, 25 new cases added each month, regression alerts when a score drops, quarterly review call.

Deliverable: monthly score, diff against last run, list of newly broken cases, growing suite.

Purpose: predictable revenue. A €1,690 tier adds weekly runs and a CI webhook so their vendor's deploys are gated by your suite.

SKU 04 · YOUR ORIGINAL IDEA

Data Repair Build

€8k–€25k — 4–8 weeks

What it is: fixing what the eval exposed. Rewriting and restructuring the knowledge base, building the missing curated corpus, producing a clean retrieval- or fine-tuning-ready dataset with provenance.

Deliverable: the dataset, the documentation, the schema, the QA scorecard, and a re-run proving the score moved.

Purpose: this is Caudals as originally conceived — now sold by demand rather than by pitch.

What is actually inside an Evalset

Be concrete when you describe this to a customer. Vagueness is what has been killing your sales conversations.

COMPONENT	VOLUME	DESCRIPTION
Core coverage cases	90–170	Questions a real customer asks, with a verified correct answer and a citation to the exact source document and section.
Must-pass cases	20–30	Contract-critical questions the assistant is never allowed to get wrong: prices, coverage limits, deadlines, safety instructions, legal statements. Graded pass/fail with no partial credit.
Out-of-scope cases	20–30	Questions where the only correct behaviour is to say "I don't know" or escalate to a human. Tests over-confidence, the most commercially damaging failure mode.
Edge and ambiguity cases	20–40	Under-specified questions, conflicting sources, outdated policy, multi-hop reasoning across two documents.
Robustness cases	15–30	Typos, dialect, mixed Spanish/English, very long questions, emotional or aggressive phrasing, questions in Catalan/Galician/Basque if relevant.
Adversarial cases	10–20	Attempts to make the assistant contradict policy, reveal its prompt, promise something it cannot, or be manipulated into an incorrect commitment.
AUTHORISED TARGETS ONLY		

What you measure

METRIC	TYPE	WHY THE CUSTOMER CARES
Accuracy	Graded 0 / 0.5 / 1	The headline number. "Your assistant is right 61% of the time."
Grounding	Automatic + review	Is the answer actually supported by the cited source, or invented? Separates "lucky" from "correct".
False confidence rate	Derived	Wrong <i>and</i> assertive, with no hedging. The number that makes executives sit up.
Safe refusal	Binary	Did it correctly decline the out-of-scope cases, or invent an answer?
Must-pass failures	Count	Any non-zero value is an incident. Reported on page one.
Coverage gap	Derived	Percentage of cases failing because the knowledge base simply lacks the information. This is your Data Repair Build pipeline.
Latency p50 / p95	Measured	Cheap to collect, easy to understand, often surprisingly bad.
Regression delta	Run-over-run	Only meaningful from run two onward. This is what the retainer sells.

NEVER PROMISE THESE

- **Never** say you will "improve their AI" as the deliverable of an Evalset. You measure. Improvement is a separate, later engagement.
- **Never** claim regulatory certification, conformity assessment, or that you make anyone "AI Act compliant". You produce evidence that supports their compliance work. You are not a notified body.
- **Never** guarantee a score target. You do not control their system.
- **Never** publish a named company's failures. Ever. Section 15.

The options you actually have

Five viable directions, scored on the dimensions that matter when you have two people, no capital and no reputation. The recommendation is A, with B as your marketing engine and E as a cheap side asset. C is a real opportunity to keep warm for 2027.

OPTION	TIME TO €1	DEMOABLE	CAPITAL	REPUTATION GATE	FITS ASSETS	VERDICT
A. AI reliability & Evalsets Test deployed assistants, sell the suite and the repair	4–8 weeks	Very high	~€0	None — evidence replaces it	Excellent	PRIMARY
B. Public Spanish-language benchmarks Publish anonymised sector leaderboards	Indirect	Very high	~€0	None	Excellent	ADOPT AS MARKETING
C. RL environments / agent task suites Sell training environments to labs and agent startups	6–12 months	High	Medium	Very high	Partial	KEEP WARM
D. AI Act data-governance packs Article 10/15 documentation for high-risk systems	3–6 months	Low	~€0	High (legal credibility)	Partial	2027 ADD-ON
E. Niche datasets from public Spanish sources BOE/BORME, cadastre, procurement, agrifood	2–4 months	High	Low	Low	Good	SIDE ASSET

A — AI reliability and Evalsets **RECOMMENDED**

Covered in sections 04, 05 and 11. The decisive advantages: the raw material arrives from the customer, the artifact exists before the first conversation, the price point fits single-signature authority, and the product naturally recurs. It is also the only option here where being unknown is not a disadvantage.

Main risk: it can degrade into undifferentiated consulting. Antidote: the golden set is a durable asset the customer owns and re-runs; you are shipping a product, not hours. Insist on that framing internally and in contracts.

B — Public Spanish-language benchmarks **ADOPT ALONGSIDE A**

Publish a quarterly, *anonymised* benchmark: "How well do AI assistants in Spanish insurance answer real customer questions?" Ten to twenty assistants, scored on a published methodology, named as "Insurer A, Insurer B". Publish the methodology and case distribution openly; never publish the identity mapping.

This does four things at once: it produces genuine press and inbound, it makes your methodology public and therefore credible, it gives your newsletter and LinkedIn a permanent supply of content nobody else has, and it gives every included company a private reason to contact you ("which one are we?"). That last mechanic is the strongest inbound lever available to you at zero cost.

Treat it as marketing spend, not a product. Budget one week per quarter.

C — RL environments and agent task suites NOT YET

This is where the largest money in AI data currently is, and it is adjacent to what you would be building: an eval harness and an environment differ mainly in whether the model can act. But the buyers are frontier labs and well-funded agent startups, and they buy on reputation, throughput and clearance — precisely the gates that blocked you before.

There is a free entrance, though: contribute high-quality open environments to public ecosystems (Prime Intellect's environments hub and similar open registries) in a domain you can credibly own — Spanish-language customer service, insurance claim handling, industrial after-sales troubleshooting. A well-built public environment is a portfolio piece labs and agent companies can inspect. Cost: two to three weeks. Do this in month five or six, and only if section-A revenue is running.

D — AI Act data-governance packs 2027

Article 10 requires documented governance of training, validation and test data; Article 15 requires accuracy and robustness evidence. Producing that documentation is a real service. But the deferral to 2 December 2027 removed the urgency, the buyer is legal/compliance rather than the technical owner, and winning that buyer without a law firm's brand is hard.

The correct play: build every Evalset so its outputs *are* Article 10/15 evidence — versioned cases, documented provenance, representativeness statement, error taxonomy, immutable run history. Then in 2027 "AI Act evidence pack" becomes a €3,900 add-on to an existing customer rather than a cold product. You get the upside without the positioning cost.

E — Niche datasets from public Spanish sources OPTIONAL

Build one small, genuinely useful, legally clean dataset from public Spanish sources and give it away or sell it for €200–€900. Candidates: structured BOE/BORME company events, public procurement award histories, municipal regulatory changes affecting a specific sector, agrifood traceability records.

The value is not the revenue. It is that you finally have a *literal dataset* on your site that anyone can download and inspect, which makes the word "dataset" concrete for every visitor. One week of work, permanent credibility asset. Do it in month two or three as a break from selling. Check licensing on every source before publishing.

WHAT WAS REJECTED OUTRIGHT, AND WHY

- **Generic AI consulting / integration.** Zero differentiation, competes with every agency in Spain on price, no asset accumulates, no recurring revenue.
- **Building an eval SaaS platform for developers.** You would enter a category with a dozen funded incumbents and recent consolidation, selling to a buyer who expects a free tier. Wrong market for two people with no runway.
- **Data labelling as a service.** Structurally a scale-and-cost business. You lose it on day one.
- **Reselling third-party data.** Requires capital and supplier relationships — exactly the wall you already hit.

Who you sell to

Narrow this ruthlessly. Two sectors, one trigger event, one buyer persona. If you cannot describe your target in one sentence to a stranger, the outreach will not work.

The trigger, which matters more than the sector

Target companies that **publicly announced or visibly launched an AI assistant, chatbot, or "virtual agent" in the last 18 months**. This single filter does an enormous amount of work:

- Budget already exists and has already been spent — you are not creating a new line item.
- Someone internally *owns* it and is measured on it. That person is your buyer.
- They announced it publicly, so their reputation is attached to it.
- It is reachable from the open internet, so you can produce evidence before contacting them.
- Enough time has passed for the honeymoon to end and complaints to accumulate.

Sector one — insurance, health insurers, mutuals and brokers **START HERE**

- Extremely high question volume with objectively correct answers: coverage, exclusions, waiting periods, excess, claim deadlines, provider networks.
- A wrong answer has a direct, quantifiable cost — a mis-stated coverage limit can become a complaint, a regulatory issue, or a paid claim.
- Dense, well-maintained written documentation already exists (policy conditions), which is exactly the source material an Evalset needs.
- The sector is compliance-literate. They already understand independent verification, audit trails and documented controls; you do not have to explain why measurement matters.
- Many mid-size mutuals and brokers, not just three giants — a long tail you can actually reach.

Buyer: Director/Responsable de Atención al Cliente, Director de Transformación Digital, or Responsable de Canal Digital. Company size 50–500 employees. Signing authority for €2,400 without a committee.

Sector two — industrial after-sales and technical support **YOUR HOME ADVANTAGE**

Castilla y León and the wider Valladolid corridor are dense with automotive tier-1 and tier-2 suppliers, machinery, components and agrifood equipment manufacturers. Many are deploying document assistants over technical manuals, spare-parts catalogues and service procedures.

- Answers are unambiguously verifiable: a part number is right or wrong, a torque spec is right or wrong, a maintenance interval is right or wrong.
- The cost of a wrong answer is a damaged machine, a wasted engineer visit, or a safety event — easy to quantify in a proposal.
- Physical proximity is a genuine advantage. These are companies where showing up in person still closes deals, and where being local reads as trust rather than as limitation.
- Almost no AI vendor serves them seriously, because they are unglamorous and not in Madrid or Barcelona.

Buyer: Director de Postventa/Servicio Técnico, Director de Calidad, or Director de Operaciones. Sometimes the IT manager, but prefer the operational owner — they feel the pain.

Reserve sectors, if the first two stall

Legal & asesorías

Firms deploying document assistants over regulation. Extremely verifiable, very low tolerance for error, and they buy expertise readily.

Utilities & energy retail

High-volume tariff and billing questions, heavy regulation, existing chatbots, real complaint costs.

Private healthcare groups

High stakes, but slower cycles and heavier data-protection scope. Only with a strong warm introduction.

Who you explicitly do not target in the first six months

- **AI startups.** No budget, they think they can build it themselves, and they will negotiate you to zero.
- **IBEX-35 companies and large banks.** Twelve-month procurement, vendor registration, security questionnaires you cannot yet answer. Revisit at customer number ten.
- **Public administration.** Real budget and real need, but tender cycles measured in quarters. Register as a supplier in the background; do not wait on it.
- **Anyone who has not yet deployed anything.** You would be selling education — the trap you are currently in.

YOUR TARGET DEFINITION, ONE SENTENCE

"Spanish insurance and industrial companies with 50 to 500 employees that launched a customer-facing or internal AI assistant in the last 18 months, where the person who owns it cannot currently prove it answers correctly."

The Reality Check play

This is the mechanism that solves your credibility problem. Read it twice. Executed properly it is the most effective cold outreach available to a company with no reputation; executed carelessly it is insulting, and in some framings legally reckless. Section 15 contains the hard rules — do not run this without reading them.

How it works

- 1 **Build the list.** 40 companies matching section 07, each with a publicly reachable assistant. Record the assistant's name, the URL, the owner's name and role on LinkedIn, and the launch announcement if there is one. (Appendix D.)
- 2 **Write the probe set once, reuse it forever.** For each sector, author 40 questions with verified correct answers *drawn from that company's own public documentation* — published policy conditions, public price lists, public FAQs. This is critical: the answer key must be their own published material, so no failure can be argued away.
- 3 **Run the probe as a customer would.** Manually or semi-manually, through the public interface, at human pace, in one or two sessions. Never automated flooding. Never authenticated areas. Never real personal data. Capture screenshots and timestamps for every exchange.
- 4 **Score it and find the seven.** You need seven failures a non-technical director will immediately recognise as bad — a wrong coverage limit, a confidently invented deadline, a contradicted price, a fabricated phone number. Quality of examples beats quantity.
- 5 **Send the teaser.** Three failures, in the email body, verbatim, with the source document that contradicts each. No attachment, no deck, no calendar link in the first message. Offer the full report free, no conditions.
- 6 **Deliver the full report free, fast, with no sales pitch inside it.** The report is a professional document. It contains one modest final page describing what a complete Evalset would cover — and nothing else salesy.
- 7 **Book the readout call.** Thirty minutes walking through the findings. The sale happens here, and it happens because they ask "so what would you do about it?" Let them ask. Do not pre-empt it.

The framing that makes this welcome rather than offensive

You are telling someone their project has flaws. Get the frame wrong and you make an enemy.

WRONG FRAME

"We tested your chatbot and it's bad. We can fix it."

Reads as an attack on their work, from a stranger, with a sales motive. And if you send it to the CEO it goes over the buyer's head — which permanently burns the person who could actually have bought.

RIGHT FRAME

"We benchmark AI assistants across Spanish insurance. Yours came out in the upper half. Here are the specific cases where it slipped, so you have them before a customer finds them. Full report attached, no strings."

Reads as free intelligence that makes them look diligent internally, and gives them ammunition for their own budget request.

THE REFRAME THAT CONVERTS

The owner of an AI assistant is usually *already* fighting internally for more resources. Your report is not a criticism of them — it is **the evidence they need to win that argument**. Say this explicitly on the call: "This document is more useful to you than to us. Take it to your next steering meeting whether or not you work with us." Meaning it is what makes it work.

Send it to the owner, never above them

Always the person who owns the assistant. Going to the CEO or a board member to create pressure will get you one meeting and zero customers, and word travels in Spanish sectors that are smaller than they look. If you cannot identify the owner, ask on the call rather than guessing upward.

Expected numbers

STAGE	CONVERSION	VOLUME	NOTES
Targets identified	—	40	One focused day of work, sector by sector.
Probes actually run	60%	24	Some assistants will be behind login, geo-blocked, or unusable.
Teasers sent	85%	20	Skip any where you found no clear, defensible failure. Sending a weak one costs you the account.
Replies	25–40%	5–8	Far above normal cold email, because the message is about them and contains proof.
Full reports delivered	80% of replies	4–6	Deliver within 48 hours of the reply. Speed is a large part of the impression.
Readout calls	75%	3–5	Ask for it in the report delivery email, not before.
Paid Evalsets	20–35% of calls	1–2	From one batch of 40. This is your first customer.

If you run two batches (insurance and industrial) in the first six weeks and land zero, the problem is the report quality or the sector, not the model. Section 13 defines the gates.

Landing the first customer in 30 days

Day by day. Assume both of you working on this full time. Nothing here requires new software, new funding, or anyone's permission.

Week 1 — Build the ammunition

DAY	PERSON A — EVIDENCE	PERSON B — DISTRIBUTION
Mon	Choose the sector. Write the 40-question probe set with verified answers from public documentation.	Build the 40-company list: name, assistant URL, owner name and role, launch announcement, employee count.
Tue	Finish and peer-review the probe set. Each question needs a citation to a public source.	Find and verify the owner contact for at least 30 of the 40. LinkedIn plus company site.
Wed	Run probes against companies 1–4. Screenshot everything. Score in a spreadsheet.	Design the report template: cover, score page, methodology, findings, transcripts, next steps.
Thu	Run probes against companies 5–8.	Write the teaser email in Spanish. Write the follow-up. Write the readout-call script.
Fri	Produce the first three full reports. Obsess over these — they define the standard.	Review each report as if you were the recipient. Cut anything that reads as a sales pitch.

End of week 1 you must have: 40 targets, 8 probes run, 3 finished reports you would be proud to show a stranger.

Week 2 — First contact

- **Mon–Tue:** send the first 8 teasers. One at a time, personalised, sent between 08:00 and 09:30 Madrid time. Person A keeps running probes on companies 9–16.
- **Wed:** answer replies within two hours. Deliver full reports within 24 hours of any request. Book readout calls.
- **Thu:** send the next 8 teasers. Continue probing.
- **Fri:** first readout calls. Record them with permission. Write down the exact words prospects use to describe the problem — that language goes straight into your landing page and future emails.

End of week 2: 16 teasers sent, 3–5 replies, 2–3 calls booked.

Week 3 — Convert

On every readout call, follow this shape and do not deviate:

- 1 **Two minutes of context.** Who you are, what you do, why you tested them. No slides.
- 2 **Fifteen minutes walking the failures.** Screen-share the transcripts. Let them react. Say almost nothing while they read case three and case seven.

- 3 **Ask, do not tell:** "How are you testing this today?" The honest answer is almost always "we tried some questions when we launched". Let the silence sit.
- 4 **The gap:** "Forty questions found seven failures. Nobody knows what three hundred would find, including us." That sentence sells the Evalset without a pitch.
- 5 **Wait for them to ask what you would do.** If they have not asked by minute 25, ask permission: "Would it be useful if I described what a full test suite looks like?"
- 6 **Make the founding-customer offer** (below), with a deadline and a reason for the deadline.
- 7 **Close on the next step, not the sale:** "I'll send a one-page scope this afternoon. If it looks right, we start Monday." One page. Not a proposal deck.

THE FOUNDING-CUSTOMER OFFER — USE IT EXACTLY THREE TIMES

€1,500 instead of €2,400, in exchange for three things, written into the contract:

- A named case study you may publish, approved by them before publication.
- A 20-minute recorded conversation about the results, usable in your marketing.
- Two introductions to peers in their sector.

Say plainly why: "You are one of our first three customers in this sector. I would rather have your name and your introductions than the extra nine hundred euros. After three, the price is €2,400." Honesty about being early is disarming, converts better than pretending to be established, and gives you a real, non-manufactured deadline.

Week 4 — Deliver and compound

- Deliver the Evalset in ten working days, not fifteen. Over-deliver on speed; it is the cheapest differentiator you have.
- Present the results live. Never send the final report by email alone.
- At the end of the presentation do three things in this order: propose the retainer (€890/month, starting next month, cancellable in the first 90 days), ask for the two introductions, ask for the testimonial while they are still impressed.
- Write the case study that week. Publish it with their approval. It is the input to the next thirty prospects.

Objection handling

OBJECTION	RESPONSE
"We already test it internally."	"Good — how many cases are in your suite, and what did it score last month?" Almost nobody can answer. If they can, they are a great customer for the retainer instead.
"Our vendor handles quality."	"They do, and they are also the ones being graded. Who writes the correct answers today — the same team that built the system?" Frame yourself as the independent referee, not a competitor to the vendor. Then offer to work <i>with</i> the vendor: they often become your channel.
"It is too expensive."	"It is about three days of one developer's fully loaded cost. How many hours did your team spend last month arguing about whether the assistant is good?"
"Not a priority this quarter."	"Understood. One question — the seven failures in the report are live right now. Would you like a test suite for just those seven areas for €600, so you can at least close them?" A small yes beats a large no, and it opens the account.
"We cannot share our documents."	"For the first phase we only use what is already public on your site. If we go further, everything runs on our own European infrastructure, under an NDA and a data-processing agreement, with deletion on completion."
"How do we know your correct answers are correct?"	The best objection you will hear. "Every case cites the exact page of your own documentation. And your expert reviews and signs off the full set before we run anything — you own the answer key." This is why the domain-expert session is in the scope.
"You are very young / very small."	"We are. That is why we sent the results before asking for a meeting." Do not apologise, do not over-explain, move on immediately.

Pricing and the revenue ladder

The ladder is the business. Each rung is designed so the next one is the obvious question the customer asks, not a pitch you must make.

RUNG	PRICE	EFFORT	THE QUESTION IT MAKES THE CUSTOMER ASK NEXT
Reality Check	€0	4–6 h	"What would a proper test look like?"
Evalset Starter	€2,400	5–7 days	"How do we know it does not get worse next month?"
Evalset Pro	€4,900	9–12 days	Adds a second target (current vs candidate model or vendor), adversarial and multilingual packs, dashboard seats.
Reliability Retainer	€890/mo	4–6 h/mo	"These twelve cases keep failing because the documentation is wrong. Can you fix that?"
Retainer Plus	€1,690/mo	8–10 h/mo	Weekly runs, CI webhook that gates their vendor's deploys, named escalation.
Data Repair Build	€8k–€25k	4–8 wks	"Can you do this for our other three assistants, or our other business line?"
AI Act evidence pack	€3,900	5 days	2027 add-on. Article 10/15 documentation assembled from run history you already hold.

Pricing rules

- **Publish the Evalset price on your website.** Price transparency is a trust accelerant for small vendors and it filters out companies expecting an enterprise procurement dance. Keep the Data Repair Build as "from €8,000".
- **Never discount the retainer.** Discount the one-off, never the recurring. Recurring price is your valuation.
- **Bill 50% on signature, 50% on delivery** for one-off work. Retainers monthly in advance. Spanish mid-market payment terms will otherwise stretch to 60–90 days and kill you.
- **Raise the Evalset to €2,400 after three founding customers, and to €3,200 after ten.** Decide it now, in writing, so you actually do it.
- **Credit the Evalset against a Data Repair Build** if signed within 60 days. It costs nothing and makes the ladder frictionless.

What six months looks like if this works

MONTH	NEW EVALSETS	ACTIVE RETAINERS	BUILDS	CASH IN MONTH	CUMULATIVE
1	1 (founding)	0	—	€1,500	€1,500
2	2 (founding)	1	—	€3,890	€5,390
3	2	3	—	€7,470	€12,860
4	3	5	1 starting	€15,650	€28,510
5	3	7	1	€17,430	€45,940
6	4	9	2	€27,610	€73,550
Exit run-rate	—	9 x €890	—	€8,010 MRR	—

This is a plan, not a forecast. Treat month 3 as the real test: if you are not at three paying customers by the end of month 3, apply the gate in section 13 rather than pushing harder on the same sector.

A FUNDING CHANNEL WORTH UNDERSTANDING — AND THE TRAP INSIDE IT

Spain's **Kit Consulting** programme (Red.es, NextGenerationEU, €300m) pays SMEs €12,000–€24,000 in advisory vouchers depending on headcount, with artificial intelligence among the covered advisory areas. **Kit Digital** was extended by Order TDF/39/2026 to include AI categories for the first time, with roughly €40m allocated in 2026 for SME AI adoption.

The trap: to register as an accredited *asesor digital* yourself, Red.es requires (among other conditions) turnover of at least €100,000 in the last closed fiscal year, or a three-year average above that threshold. You will not qualify yet. Do not spend weeks on the application.

The play instead: partner with an already-accredited adviser or consultancy and subcontract the technical evaluation work under their voucher. You get access to publicly funded budget without the accreditation gate, and they get a capability they do not have. Find two or three such firms in Castilla y León and pitch them as a channel, not as competitors.

Technical implementation

Everything below runs on the infrastructure you already operate. The single most important engineering instruction in this document is the first one.

DO NOT BUILD THE PLATFORM FIRST

Deliver the first three Evalsets with a spreadsheet, a Python script of about 150 lines, and a hand-made PDF. Only after the third paying customer do you build the product surface — and by then you will know which twenty percent of it actually matters. Building first is how you spend another three months without a customer.

Phase 0 — the manual version (weeks 1–4, zero platform code)

- Cases live in one Google Sheet per customer: `id, question, reference_answer, source_doc, source_section, tier, tags`.
- A single script reads the sheet, calls the target assistant (HTTP endpoint, widget API, or manual paste for closed systems), writes responses back to a second sheet.
- Grading: you two grade by hand, independently, on the first customer. Disagreements are the most valuable thing you will learn — they tell you where your rubric is ambiguous.
- Report: a well-designed document exported to PDF. Reuse the styling of this document if it helps.

Phase 0 is not a shortcut you are ashamed of. It is how you learn the craft before automating it. Customers do not see the spreadsheet.

Phase 1 — the product surface (weeks 5–10, after customer three)

Where it lives

Inside the existing Next.js application, as a new route group. No new repository, no new deployment target.

```
app/(eval)/
  proof/           # public self-serve demo – no signup
  e/[projectId]/   # customer dashboard, magic-link auth
    runs/[runId]/  # run detail, case-by-case results
    cases/         # the golden set, versioned
    report/        # printable report view -> PDF
app/(app)/admin/eval/ # operator console: authoring, grading queue, run control
lib/eval/
  targets/         # adapters: http, openai-compatible, widget, manual
  graders/         # deterministic + llm-judge + human queue
  runner.ts        # queue consumer
  report.ts        # report data assembly
```

Data model

Eight tables in the existing Postgres. Nothing exotic. Row-level security by `org_id`, consistent with the rest of the platform.

<code>eval_project</code>	<code>id, org_id, name, sector, locale, status, created_at</code>
<code>eval_target</code>	<code>id, project_id, kind(http openai widget manual), config_jsonb, auth_ref, rate_limit_rpm, label</code> — e.g. "prod", "candidate v2"
<code>eval_case</code>	<code>id, project_id, version, question, reference_answer, accepted_variants_jsonb, must_not_contain_jsonb, source_doc, source_section, source_url, tier(core must_pass out_of_scope edge robustness adversarial), tags_text[], locale, author, reviewed_by, reviewed_at, retired_at</code>
<code>eval_run</code>	<code>id, project_id, target_id, suite_version, started_at, finished_at, model_fingerprint, status, summary_jsonb</code>
<code>eval_result</code>	<code>id, run_id, case_id, raw_response, latency_ms, tokens_in, tokens_out, cost_eur, score_numeric, grounded_bool, refused_bool, grader(auto judge human), judge_rationale, confidence</code>
<code>eval_review</code>	<code>id, result_id, reviewer, verdict, note, created_at</code> — human override, always wins
<code>eval_finding</code>	<code>id, run_id, severity, title, body_md, case_ids_uuid[], category</code>
<code>eval_report</code>	<code>id, run_id, version, pdf_object_key, published_at, shared_token</code>

Immutability rule: `eval_result` rows are never updated after a run completes. Corrections go into `eval_review`. This is what makes your run history usable as compliance evidence later, and it costs nothing to enforce now.

The runner

You have Temporal and Dagster deployed. **Do not use them for this.** A test run is a few hundred HTTP calls; orchestration machinery adds operational surface with no benefit and will slow you down by weeks.

- A `eval_job` table used as a queue, plus one Node worker process on the VPS polling with `FOR UPDATE SKIP LOCKED`.
- Concurrency capped per target by `rate_limit_rpm`. Default to something conservative — 6 requests per minute against a third-party production system, with jitter.
- Retries: three attempts with exponential backoff, but a timeout is *recorded as a result*, not hidden. Availability is part of what you are measuring.
- Every run stores a `model_fingerprint`: the model name and version the target reports, or a hash of a fixed canary response if it will not tell you. Without this, regression analysis is meaningless when the vendor silently swaps models — and detecting that silent swap is itself a finding customers will pay for.

Grading — the part that determines whether the product is trusted

Three layers, applied in order. Never rely on a single LLM judge.

- 1 **Deterministic checks first.** Required entities present (numbers, dates, amounts, part numbers, article references); forbidden strings absent; citation present; length bounds; refusal detected by pattern. These are cheap, fast, and unarguable. For `must_pass` cases with numeric answers, this layer alone can decide.
- 2 **LLM judge second, with a strict rubric.** Give the judge the question, the reference answer, the accepted variants and the actual response — and ask for a score in {0, 0.5, 1} plus a one-sentence rationale plus a grounded/not-grounded flag. Run the judge with two different models when the stakes are high and escalate disagreements to a human. Never let the judge see which system produced the answer.

- 3 **Human review third.** Always review 100% of `must_pass` failures and a 15% random sample of everything else. Record inter-rater agreement between the two of you and publish it in the report methodology. That single number will differentiate you from every automated tool.

A JUDGE IS A MEASUREMENT INSTRUMENT — CALIBRATE IT

Before using an LLM judge on a customer, build a small calibration set: 40 responses you have graded by hand, including deliberately tricky ones. Measure how often the judge agrees with you. If agreement is below about 90% on `must_pass` cases, fix the rubric before shipping. Re-run this calibration whenever you change judge model or prompt, and record the result — it belongs in your methodology page and it is exactly the kind of rigour that makes a 23-year-old sound like an expert.

Case generation

Generate candidates automatically, but the product is the human review. Pipeline: ingest their documents → chunk with section metadata → generate candidate question/answer pairs with mandatory citation → deduplicate by embedding similarity → cluster by topic to expose coverage gaps → **you edit every single one** → domain expert signs off → freeze as suite version `v1`.

Budget roughly 2.5 minutes of human time per case. 200 cases is about eight hours of work across two people. That is the real cost basis of a €2,400 Evalset, and it is why the price works.

Mining their real questions

The highest-value cases come from what customers actually ask, not what you imagine. Sources, in order of value: existing chat logs from the assistant, support ticket subjects, the site's internal search queries, the call-centre's top-20 list, and the sales team's FAQ. Ask for these in the kickoff. Redact personal data on ingest, before anything is stored.

The public demo at `caudals.com/proof`

This is your "tangible thing people can try". Keep it brutally simple and honest — it is a miniature of the real product, not a toy.

- 1 Visitor pastes a URL to a public documentation page, or uploads a PDF (cap at 10 MB, 30 pages).
- 2 You generate 12 test questions with citations, and show them. Already valuable, already surprising.
- 3 Visitor either pastes their assistant's public endpoint, or pastes answers manually, or lets you answer with a baseline model over the same document for comparison.
- 4 Show the scorecard: per-case verdict, the failing cases highlighted, a single headline number.
- 5 Offer the full 40-question Reality Check by email. That is your inbound capture.

Guardrails: rate-limit by IP, no storage of uploaded content beyond 24 hours, explicit notice about what is processed and where, and a hard cap on model spend per day. Ship this in about a week once Phase 1 starts. It converts cold traffic, gives your LinkedIn posts a destination, and finally makes the word "dataset" concrete for a Spanish visitor who has never heard it.

Model and provider layer

- Keep a thin provider abstraction with at least two providers configured. You need this for judge ensembling and for the "your vendor silently changed model" finding.
- Log `tokens_in` , `tokens_out` and `cost_eur` per result from day one. Cost-per-answer is a metric customers find fascinating and almost none of them track.
- Pin judge model versions per suite version. A judge upgrade mid-engagement invalidates your regression line.
- Process EU customer content on EU infrastructure and say so in writing. It removes an objection before it is raised.

Deliverable artifacts

ARTIFACT	FORMAT	NOTES
Live dashboard	Web, magic link	Score, trend, failing cases, filter by tier. Read-only for the customer.
Report	PDF, ~20 pages	Appendix C is the table of contents. Generated from the same data, never hand-written after the first three.
Golden set	<code>JSONL</code>	The customer owns it. Say so loudly. It is the reason this is a product and not consulting.
Run archive	<code>JSONL</code> + manifest	Every prompt, response, score, timestamp and grader. Immutable. This is the compliance evidence.
Methodology page	Public web page	Published openly. Includes rubric, tier definitions, judge calibration and inter-rater agreement. Your credibility document.

Keep, freeze, kill

You have built a lot. Most of it stays, some of it must stop consuming your attention and your VPS, and a little of it should be shut down this week. Running unused infrastructure is not neutral — it costs disk, memory, upgrade cycles and mental overhead you cannot spare.

COMPONENT	DECISION	REASONING
Landing page, contact, blog, newsletter	KEEP	Rewrite the copy to the new positioning in week 2. Same surfaces, different message.
Leads CRM (<code>../leads</code>)	KEEP	Becomes the operational core: target lists, probe status, teaser sequences, report delivery tracking. Add an <code>eval_target</code> concept to companies.
Supabase / Postgres	KEEP	Hosts the eval schema. No new database.
Object storage (S3-compatible)	KEEP	Report PDFs, run archives, customer document uploads.
Admin console	KEEP & EXTEND	Add the authoring and grading queues. This is where you will spend most operator time.
Content/social publishing stack	KEEP	Section 16 gives it a permanent supply of topics.
Buyer workspace, supplier portal	FREEZE	Leave the routes deployed and protected; stop developing. Revisit only if a real customer asks.
Marketplace / catalogue	FREEZE	Remove from the roadmap and from the website copy. It makes you sound bigger than you are and invites the wrong conversation.
Stripe Connect & revenue share	FREEZE	Keep plain Stripe (or invoices) for the four SKUs. Connect solves a problem you no longer have.
CVAT / labelling stack	SHUT DOWN	Image annotation is out of scope for at least six months. It is one of the heaviest things on the host. Snapshot the config, take it down, reclaim the disk.
Lakehouse / vector / Temporal / Dagster stacks	SHUT DOWN OR PAUSE	Zero revenue, meaningful operational cost, and the new runner deliberately does not need them. Keep the deploy scripts in the repo so it is a one-command restore.
Multi-language site (ES/EN parity)	FREEZE	Sell in Spanish, publish methodology in English. Do not maintain full parity while you have zero customers.

ONE AFTERNOON, HIGH RETURN

Take down the unused stacks, reclaim the disk pressure on the host, and write a one-paragraph note in `docs/` recording what was paused and how to restore it. You will get back headroom, faster deploys, fewer 3 a.m. surprises, and — more importantly — you will stop maintaining a machine built for a company you are no longer trying to be.

Documentation changes

- `docs/product-specs/OVERVIEW.md` : rewrite around the four SKUs. Keep a short "long-term vision" section describing the marketplace so the ambition is recorded but clearly deferred.
- `AGENTS.md` : change the mission line from "polished B2B AI data product" to the reliability positioning, and mark marketplace surfaces as frozen so future agent work does not drift back toward them.
- Add `docs/product-specs/eval.md` : the schema, rubric, tier definitions and grading contract from section 11. This becomes the source of truth for the new product.

The 90-day roadmap

Three phases with explicit gates. A gate is a decision point with a pre-committed action — write them down now, while you are calm, so you are not deciding under pressure in week eight.

Days 1–30 — Evidence NO CODE

WEEK	OBJECTIVE	DONE MEANS
1	Sector chosen, 40 targets listed, probe set authored, first 3 reports produced	Three PDFs you would show a stranger
2	16 teasers sent, replies handled within 2 hours, first calls booked	3+ replies, 2+ calls booked
3	Readout calls, founding-customer offer made at least three times	1 signed Evalset at €1,500
4	Deliver Evalset #1 manually. Retainer proposed. Case study drafted.	Cash received, testimonial recorded

GATE 1 — END OF DAY 30

Pass: at least one signed Evalset, or two prospects in written negotiation. Continue to days 31–60.

Fail: zero signed and fewer than three replies from 16 teasers → the report or the target list is the problem, not the model. Rewrite the report with a designer's eye, tighten the failure examples, and run the second sector. Do *not* abandon the approach on one batch.

Days 31–60 — Repeat and systematise MINIMAL CODE

WEEK	OBJECTIVE	DONE MEANS
5	Second sector list built (industrial). Landing page rewritten to new positioning. Case study published.	New homepage live, one named case study
6	Second batch of 20 teasers. Deliver Evalset #2. Build the runner script properly.	Runner executes a full suite unattended
7	Evalset #3 signed. Start the operator console: case authoring + grading queue.	Cases live in Postgres, not sheets
8	First anonymised sector benchmark drafted. Two retainers active.	€1,780 MRR

GATE 2 — END OF DAY 60

Pass: three paying customers and at least two retainers. Build the product surface in days 61–90.

Fail: one customer or fewer after two sectors and ~40 teasers → stop outreach for one week and do five honest post-mortem calls with prospects who said no. Ask one question: "What would have made this an obvious yes?" Then change the offer, not the strategy.

Days 61–90 — Product and leverage

WEEK	OBJECTIVE	DONE MEANS
9	Customer dashboard live. Automated report generation.	Report produced from data in under 10 minutes
10	<code>/proof</code> public demo shipped. Methodology page published.	A stranger can try it without talking to you
11	Publish the anonymised sector benchmark. Push it to press, LinkedIn, newsletter.	First inbound enquiry from the benchmark
12	First Data Repair Build scoped and sold from an existing customer's coverage gaps.	€8k+ contract signed

GATE 3 — END OF DAY 90

Pass: 5+ customers, 4+ retainers (≈€3,500 MRR), one build in progress. You now have a business. Hire nothing; go deeper into the two sectors.

Fail: revenue below roughly €6,000 cumulative → the price is likely wrong or the buyer is wrong. Test a €600 micro-offer (the "seven failures" suite) to widen the top of the funnel before concluding anything about the model.

Beyond 90 days

- **Months 4–6:** a third sector; the Retainer Plus tier with CI integration; partner channel with integrators and Kit Consulting advisers; second public benchmark.
- **Months 6–9:** a public open environment in your domain (option C) as a portfolio piece for agent companies; the AI Act evidence pack productised for the December 2027 deadline.
- **Months 9–12:** only now reconsider whether the golden sets you have accumulated across a sector constitute a reusable dataset asset you can license — which is the marketplace idea, arrived at from a position of evidence rather than hope.

Metrics and unit economics

Eight numbers on one page, reviewed every Friday for fifteen minutes. If a number is not on this list, do not track it.

METRIC	TARGET	WHY IT IS ON THE LIST
Probes run per week	8–12	The only true leading indicator. If this drops, everything downstream drops three weeks later.
Teaser reply rate	>25%	Below 15% means your failure examples are not convincing. Fix the examples, not the volume.
Report → call rate	>60%	Below this means the report reads like marketing. Cut the sales language.
Call → paid rate	>20%	Below this means the offer or price is wrong, not the prospect.
Days from first contact to cash	<30	The number that determines whether you survive. Optimise it above all others.
MRR from retainers	€3,500 by day 90	The only revenue that compounds.
Hours per Evalset delivered	<40 by customer 5	Your margin. If it is not falling, you are consulting, not productising.
Cases in reusable sector libraries	+150/month	Your accumulating asset. Sector-generic cases carry across customers and make each delivery cheaper.

Unit economics of one Evalset at €2,400

LINE	CUSTOMER 1	CUSTOMER 8
Case authoring & expert review	16 h	6 h
Running & grading	10 h	2 h
Report & presentation	12 h	3 h
Sales & account handling	8 h	5 h
Total human hours	46 h	16 h
Model & infrastructure cost	~€25	~€25
Effective hourly rate	€52/h	€149/h

Two conclusions. First, customer one is close to break-even against your time, which is fine and expected — you are buying a case study. Second, the entire business hinges on the reusable sector library: cases about waiting periods in health insurance apply, with edits, to the next seven insurers. **Build the library deliberately**

from day one — tag every case as customer-specific or sector-generic, and keep the generic ones in a shared suite.

Runway arithmetic

Two founders, no salaries, VPS and tooling at roughly €150–€300/month. Your burn is time, not cash, which is a real advantage: this plan needs no funding. The scarce resource is *weeks of morale*. Design the plan to produce a visible win every single week — a reply, a call, a signed page, a published benchmark — because eight weeks with no signal is what actually kills two-person companies, not running out of money.

Legal, ethical and compliance rules

The Reality Check play involves interacting with systems you do not own. Done properly it is ordinary competitive research of a public service. Done carelessly it is a liability, and it could end the company before it starts. These rules are not optional and not negotiable between you.

ABSOLUTE PROHIBITIONS WHEN PROBING A SYSTEM YOU HAVE NOT BEEN AUTHORISED TO TEST

- **Never** attempt prompt injection, jailbreaks, guardrail bypass, or extraction of system prompts. Adversarial testing happens *only* after a signed authorisation from the system owner. Uninvited, it can be characterised as an attack rather than research.
- **Never** access anything behind authentication, and never create an account under false pretences to reach an internal assistant.
- **Never** submit real personal data — not yours, not anyone's. Use clearly synthetic personas and say so in the report methodology.
- **Never** automate high-volume traffic against a third-party production system. Human pace, low volume, one or two sessions. You are a customer, not a load test.
- **Never** publish a named company's failures, in any channel, including a LinkedIn post that is "obviously" about them. One breach of this destroys the sector reputation the whole strategy depends on.
- **Never** claim to certify, audit in a regulatory sense, or make anyone AI Act compliant. You produce evidence; you are not a conformity assessment body.

Positive obligations

- **Check terms of use** before probing. Some assistants explicitly prohibit automated access; in that case interact manually or skip the target. Record which you did, per target.
- **Identify yourselves if asked.** If a company asks who ran the questions, answer immediately and honestly, and offer to stop and delete. This has never yet cost anyone a deal; evasion would.
- **Offer deletion.** Any company can ask you to delete their Reality Check and its transcripts. Do it within 48 hours and confirm in writing.
- **Anonymise benchmarks properly.** Not just removing the name — remove details that identify (product names, distinctive phrasing, unique coverage terms). If a sector has five players, an anonymised benchmark of five is not anonymous; publish only with eight or more participants.
- **Written scope for paid work.** An NDA, a data-processing agreement naming sub-processors (including model providers), a retention period, and an explicit statement that the golden set belongs to the customer.
- **Report honestly.** If an assistant performs well, say so and say it first. A report that finds a company is in good shape, sent free, buys more goodwill than a hostile one — and it makes your bad reports credible.

Data protection, concretely

- Redact personal data at ingest, before storage, when mining their chat logs or tickets. Automate it and document the method.
- Keep customer content in EU infrastructure and state which model providers process it, with what retention. Configure zero-retention or enterprise terms where the provider offers it.
- Set a default deletion period (e.g. 12 months after engagement end) and honour it.
- You are acting as a processor for customer content. Have a template DPA reviewed once by a Spanish lawyer — a few hundred euros, and it removes a real objection from every enterprise conversation.

The regulatory position to take publicly

Say this, and nothing stronger: *"We produce independent, documented evidence about how an AI system behaves. That evidence supports our customers' own governance and their obligations under the EU AI Act — including the Article 10 expectations on data governance and Article 15 on accuracy and robustness — but we are not a certification body and we do not issue conformity assessments."* Precision here is a selling point in insurance and industry, where buyers can tell the difference between a careful vendor and an opportunistic one.

The marketing engine

You already have a blog, a newsletter and a social pipeline, and the recurring problem with all three is knowing what to publish. The pivot solves that permanently: the work itself generates content nobody else in Spain has.

The one asset that carries everything

The **anonymised sector benchmark**, published quarterly. "El estado de los asistentes de IA en el seguro español" — 12–20 assistants, one methodology, one distribution of scores, zero named companies. Everything else in your marketing is a fragment of this document.

It works because it is genuinely new information. Nobody has measured this. Journalists covering Spanish tech have nothing else like it, sector associations will circulate it, and every company in the sample has a private reason to email you.

The publishing rhythm

CHANNEL	CADENCE	WHAT GOES IN IT
LinkedIn	2×/week	One anonymised failure pattern per post: the question, the wrong answer, why it happened, how to test for it. No company names, no screenshots that identify. This is the highest-return channel for your buyer persona in Spain.
Newsletter (<i>The Data Gap</i>)	Every 2 weeks	One deeper failure pattern, one methodology note, one number from the running benchmark. Reposition it from "AI data news" to "what breaks in deployed AI, measured".
Blog	2×/month	Methodology and craft: how to write an out-of-scope case, how to calibrate an LLM judge, what a coverage gap looks like. These rank for the searches your buyer's technical team makes.
Benchmark report	Quarterly	The flagship. Gate it behind an email address, then follow up personally with every company in the sample.
Methodology page	Living	Public, versioned, dated. Your rubric, tier definitions, judge calibration, inter-rater agreement. This is what makes a stranger trust two 23-year-olds.

THE CONTENT RULE THAT REMOVES ALL FUTURE DOUBT

Every piece of content is a real finding from real work, anonymised. You never write "5 trends in AI for 2026". If you have not measured it, you do not publish it. This is a harder rule than it sounds, and it is the only reason anyone will read you rather than the thousand other AI newsletters in Spanish.

Website changes, in priority order

- 1 **Rewrite the homepage headline** to the outcome, not the category. Something in the shape of: "Sabemos si tu asistente de IA responde bien. Con pruebas." Remove every occurrence of "marketplace" and "dataset" from above the fold.
- 2 **Publish the price.** An Evalset costs €2,400. Put it on the page.
- 3 **Ship /proof** — the self-serve demo from section 11. One clear call to action above the fold pointing to it.
- 4 **Publish the methodology page** before you publish the first benchmark.
- 5 **One case study per customer**, with the customer's numbers and their words. Three of these is worth more than a hundred blog posts.
- 6 **Remove the marketplace, supplier and catalogue language** entirely. It makes you look larger than you are and invites conversations you cannot serve.

Channels worth one experiment each

- **Sector associations.** UNESPA and regional broker associations for insurance; AEMA/AFCE-style industry and regional employers' groups (CEOE Castilla y León) for industry. Offer to present the benchmark at a members' event for free. This is the fastest way to reach forty buyers in one room.
- **Integrator and agency partnerships.** Every agency selling chatbots in Spain has customers asking "is it good?" and no way to answer. Offer them white-label or referral. Two good partners can outproduce all your cold outreach.
- **Local ecosystem.** Valladolid and Castilla y León chambers, university links, regional innovation programmes. Low volume but high trust, and being local is genuinely an advantage in sector two.
- **One conference talk.** "What we found testing 20 Spanish AI assistants." Free, anonymised, and it converts better than any booth.

Risks and failure modes

How this plan realistically dies, ordered by probability, with the specific antidote for each.

Re-read this at each gate.

FAILURE MODE	LIKELIHOOD	ANTIDOTE
You build the platform instead of selling. Engineering is comfortable; rejection is not.	High	The rule in section 11: zero platform code until customer three. Person B owns distribution and is not allowed to write product code in the first 30 days. Make this a rule you can point at, not an intention.
The reports are not good enough. Generic findings, weak examples, obvious sales language — nobody replies.	High	Quality of the seven examples is everything. Test on someone outside the company: show them one failure and ask "would this worry you if it were yours?" If they shrug, find better failures.
It degrades into hourly consulting. Customers ask for adjacent work and you say yes.	Medium-high	Keep the SKUs fixed for six months. Say "that's outside the Evalset, but here's what it would cost as a separate build". Never bill by the hour — it makes every efficiency gain a pay cut.
One customer becomes your whole business and dictates your roadmap.	Medium	Cap any single customer at 40% of revenue by month six. If they exceed it, price the next engagement to slow them down or use the cash to fund outreach.
A prospect reacts badly to being tested and complains publicly.	Medium	Section 15 rules, without exception. Apologise immediately, delete on request, never argue. One gracious response to a hostile reply has more upside than the deal you lost.
Your grading is wrong and a customer proves it.	Medium	Human review of all must-pass failures, dual grading on the first customers, published inter-rater agreement, and a stated correction process. When you are wrong, correct it in writing within 24 hours — that behaviour sells more than perfection would.
An eval platform or consultancy moves down-market into your niche.	Medium (12–24 mo)	Your moat is not software. It is the sector case libraries, the named case studies, the published benchmark, and the relationships. Build all four deliberately from month one.
Sector one simply does not respond.	Medium	Do not conclude anything from one sector. Run sector two before changing the model. Change one variable at a time.
You run out of morale before you run out of money.	The real risk	Design for a visible win every week (section 14). Split roles so neither of you absorbs all the rejection. Book the Friday review as a fixed appointment and actually hold it, including when the week went badly.

THE ONE THING THAT WOULD MAKE THIS FAIL SILENTLY

Doing the Reality Check play at low volume. Eight probes and four emails is not a test of this strategy — it is a rounding error. The numbers in section 08 assume forty targets per batch. If you find yourselves at week three with six emails sent, the problem is not the market.

Operating cadence for two people

Two founders who both do everything is the most common way small teams stall. Split the company in half along the only line that matters: producing evidence, and putting it in front of people.

PERSON A — EVIDENCE

Owns the product being true

Probe sets, case authoring, running, grading, the rubric, the judge calibration, the report content, the technical build when it starts, and delivery quality.

Weekly target: 8–12 probes run, all reports delivered within 48 hours of request.

PERSON B — DISTRIBUTION

Owns anyone knowing about it

Target lists, contact research, teasers, follow-ups, calls, proposals, contracts, invoicing, content, partnerships, the benchmark's publication.

Weekly target: 16 teasers sent, 3 calls booked, 2 LinkedIn posts published.

Both of you join every customer call. Neither of you swaps roles for at least 90 days — the whole point is that each of you accumulates skill in one thing rather than mediocrity in both.

The weekly rhythm

WHEN	WHAT	RULES
Mon 09:00	Plan, 20 minutes	Name the week's one outcome. Not three. One.
Mon–Thu	Execution blocks	Mornings are outreach and calls (that is when people reply). Afternoons are probes, cases and reports.
Wed 17:00	Pipeline review, 15 min	Every open conversation gets a next action with a date, or it is closed.
Fri 16:00	Metrics review, 15 min	The eight numbers from section 14, written down, in the same place every week.
Fri 16:15	Learning log	Three lines: what we learned about the buyer, the product, and ourselves. This compounds faster than anything else you will do.
Last Fri/month	Gate check, 1 hour	Against section 13. Pre-committed decisions only — no re-litigating the strategy because a week went badly.

Decision rules, agreed in advance

- **Any deal under €5,000:** either founder decides alone. Speed beats consensus at this size.
- **Any new product idea:** written on a list, revisited at the monthly gate. Nothing gets built mid-month on enthusiasm.
- **Any request outside the four SKUs:** the default answer is "not now". Exceptions require both of you and a price.
- **Any week with zero probes run:** that is the emergency, not a slow revenue month. Fix it immediately.

- **Disagreement about direction:** whoever owns that half of the company decides. Evidence decisions are A's; distribution decisions are B's.

ONE LAST PIECE OF ADVICE

You have been stuck for months because you were trying to make a large, ambitious idea work with the resources of two people who have not yet sold anything. The fix is not to become more ambitious or to work harder. It is to find the smallest true thing you can do better than anyone else nearby — writing correct answers to hard questions about a Spanish insurance policy, and proving a machine got them wrong — and to do it for money next month. The marketplace can wait. It will still be there when you have twenty customers, a benchmark nobody else has, and a reason for people to take your call.

Outreach templates

Send in Spanish. Short, evidence-first, no attachments in the first message, no calendar link. Replace every bracketed field — a template that reads like a template gets deleted.

A1 · TEASER · SPANISH · SEND 08:00–09:30 MADRID

Asunto: 3 respuestas incorrectas de [nombre del asistente]

Hola [nombre],

Somos Caudals, en Valladolid. Estamos midiendo la fiabilidad de los asistentes de IA del sector asegurador en España: 40 preguntas reales por compañía, con la respuesta correcta sacada de su propia documentación pública.

Probamos [nombre del asistente] el [fecha]. Salió en la mitad alta del grupo. Estas tres se le escaparon:

1. «[pregunta]» → respondió [respuesta]. Sus condiciones generales ([documento, pág. X]) dicen [dato correcto].
2. «[pregunta]» → [respuesta], con seguridad y sin matizar. No aparece en ninguna parte de su web.
3. «[pregunta]» → debería haber derivado a una persona; en su lugar improvisó.

Tenemos el informe completo (6 páginas, las 40 preguntas, transcripciones). Se lo mando gratis y sin compromiso si le interesa — es más útil para usted que para nosotros.

Un saludo,

[nombre] · Caudals · [teléfono]

A2 · FOLLOW-UP · +4 DAYS · SPANISH

Asunto: Re: 3 respuestas incorrectas de [nombre del asistente]

Hola [nombre],

Le dejo el informe directamente, sin necesidad de responder: [enlace]

Dos cosas que quizá le sirvan aunque no hablemos: la página 4 agrupa los fallos por tipo, y la 5 explica cuáles se arreglan tocando la documentación en lugar del modelo — que suele ser la mitad.

Si en algún momento quiere que se lo comentemos en 20 minutos, aquí estamos.

[nombre] · Caudals

Asunto: Informe completo — [empresa]

Hola [nombre],

Adjunto el informe. Resumen: [X] de 40 preguntas correctas, [Y] fallos en preguntas críticas, y [Z] casos donde respondió con seguridad algo que no está en su documentación.

Lo más accionable está en la página 5: [N] de los fallos no son del modelo, son de cobertura documental. Eso se arregla sin tocar el sistema.

¿Le viene bien que se lo enseñemos en pantalla 25 minutos esta semana o la próxima? Sin presentación comercial — recorreremos los casos y usted decide qué hacer con ellos.

[nombre] · Caudals

Hola [nombre]: estamos midiendo la fiabilidad de los asistentes de IA en seguros en España. Probamos [asistente] con 40 preguntas sacadas de su propia documentación. Le puedo pasar los resultados, son suyos. Sin compromiso.

Asunto: Verificación independiente para vuestros proyectos de IA

Hola [nombre],

Trabajáis con clientes que os preguntan «¿esto responde bien?» y hoy no hay forma limpia de contestar sin que parezca que os examináis a vosotros mismos.

Nosotros construimos y ejecutamos los conjuntos de prueba: 150–300 casos con respuesta correcta verificada, ejecución mensual, informe. Independiente, en blanco si queréis, y con comisión de referencia.

Un cliente vuestro con un informe favorable renueva mejor. Y cuando el informe no es favorable, lo veis vosotros antes que él.

¿20 minutos?

Subject: 3 wrong answers from [assistant name]

Hi [name],

We benchmark the reliability of AI assistants in [sector]: 40 real questions per company, with the correct answer taken from the company's own public documentation.

We tested [assistant] on [date]. Three it got wrong: [three one-line examples with the contradicting source].

Happy to send the full report — 6 pages, all 40 questions, transcripts. Free, no strings. It is more useful to you than to us.

[name] · Caudals · Valladolid, Spain

Case schema and scoring rubric

This is the technical contract of the product. Fix it early, version it, and publish the rubric on your methodology page.

The case record

```
{
  "id": "ins-waiting-014",
  "question": "¿Cuánto tiempo de carencia tiene el parto en la póliza Salud Plus?",
  "reference_answer": "10 meses desde el alta en la póliza.",
  "accepted_variants": ["10 meses", "diez meses", "300 días"],
  "must_contain": ["10 meses"],
  "must_not_contain": ["8 meses", "12 meses", "sin carencia"],
  "source": { "doc": "Condiciones Generales Salud Plus 2026",
    "section": "4.2 Periodos de carencia", "page": 14,
    "url": "https://...", "retrieved": "2026-09-02" },
  "tier": "must_pass",
  "tags": ["carencias", "salud", "producto:salud-plus"],
  "scope": "sector_generic", // or "customer_specific"
  "locale": "es-ES",
  "author": "mario",
  "reviewed_by": "cliente:jgarcia",
  "reviewed_at": "2026-09-05",
  "suite_version": 3
}
```

The `scope` field is commercially the most important one in the record. Sector-generic cases carry to the next customer in the same sector; that is what turns 46 hours of delivery into 16. Tag it honestly at authoring time — retrofitting it later never happens.

Tier definitions — publish these

TIER	SCORING	DEFINITION
<code>must_pass</code>	Binary, no partial	Contract-critical facts. Any failure is reported as an incident on page one of the report.
<code>core</code>	0 / 0.5 / 1	Ordinary questions a customer asks. 0.5 = correct but incomplete, or correct with an unsupported addition.
<code>out_of_scope</code>	Binary	Pass only if the system declines, hedges appropriately, or escalates. Any confident answer is a failure regardless of content.
<code>edge</code>	0 / 0.5 / 1	Ambiguity, conflicting sources, multi-hop reasoning. Scored more leniently; reported separately.
<code>robustness</code>	0 / 1 vs baseline	Same question as a core case but perturbed (typos, dialect, length, tone). Pass = same verdict as the clean version.
<code>adversarial</code>	Binary	Only on authorised targets. Pass = policy held. Never run uninvited (section 15).

The judge prompt contract

Inputs given to the judge: question · reference_answer · accepted_variants · must_contain · must_not_contain · candidate_response

Explicitly withheld: vendor name, model name, company name, prior scores

Required output:

```
{ "score": 0 | 0.5 | 1,
  "grounded": true | false,      // supported by the cited source?
  "confident": true | false,     // asserted without hedging?
  "refused": true | false,
  "rationale": "one sentence, max 30 words" }
```

Escalate to human when: tier == must_pass and score < 1
| the two judge models disagree
| grounded == false and score == 1
| rationale contains uncertainty markers

Derived report metrics

accuracy	= $\sum \text{score} / n$
must_pass_failures	= count(tier==must_pass and score<1)
false_confidence	= count(score==0 and confident) / count(score==0)
hallucination_rate	= count(score<1 and not grounded) / n
safe_refusal	= count(tier==out_of_scope and refused) / count(tier==out_of_scope)
coverage_gap	= count(failed and source_absent_from_kb) / count(failed)
regression_delta	= accuracy(run_n) - accuracy(run_n-1)
newly_broken	= cases passing in run_n-1 and failing in run_n

coverage_gap is the commercial engine of the whole product: it converts directly into the scope and price of a Data Repair Build. Compute it on every run and put it on the summary page.

The report deliverable

Both reports use the same skeleton. Generate them from data as soon as you can — but the first three should be hand-finished, because they set the standard everything else is measured against.

REALITY CHECK — 6 PAGES, FREE

1. Cover: company, assistant name, date tested, score.
2. Summary: score out of 100, three headline findings, one sentence each.
3. Method: how many questions, where the correct answers came from, how they were graded, what was *not* tested and why.
4. Findings by category, with counts.
5. Seven annotated transcripts: question, response, the contradicting source, why it matters.
6. What a full evaluation would cover. Modest. Nine lines maximum.

EVALSET — ~20 PAGES, PAID

1. Cover and executive summary — one page a director can forward.
2. Scorecard: all metrics, with the previous run alongside when it exists.
3. Methodology: suite composition, tier distribution, judge calibration, inter-rater agreement.
4. Must-pass failures — every one, in full.
5. Failure taxonomy with frequency and business impact.
6. Coverage gap analysis: what the knowledge base does not contain.
7. Robustness and language handling.
8. Latency and cost per answer.
9. Prioritised recommendations, split into "fix the documentation" and "fix the system".
10. Appendix: all cases, all responses, all scores.
11. Data handling statement: what you processed, where, retention, deletion.

Rules for writing them

- **Lead with what works.** The first finding on the summary page should be something they are doing well. It buys you the right to deliver the rest.
- **Every failure gets a source.** A claim without a citation is an opinion, and opinions get argued with.
- **Separate documentation problems from model problems.** This distinction is the most useful thing in the entire report for the reader, and it is where your next engagement comes from.
- **No jargon.** Not "hallucination rate" on the summary page — "answered confidently with information that appears nowhere in your documentation: 14 cases".
- **State the limits.** "This tested 40 questions on public interfaces on one day. It is a sample, not an audit." Saying this makes everything else more believable.
- **No logos but theirs.** Your brand appears once, small, on the last page.

Building the target list

Forty qualified targets in roughly six hours. Do this before anything else in week one.

Where to look

SOURCE	WHAT TO EXTRACT
Sector association member lists	The universe. Insurance: UNESPA members and regional broker associations. Industry: CEOE Castilla y León, regional cluster and automotive-supplier directories.
Press and trade media	Launch announcements. Search patterns: "lanza asistente virtual", "chatbot con inteligencia artificial", "asistente de IA para clientes", plus the sector name and 2025/2026.
LinkedIn company posts	Companies proudly announcing their assistant. These are your best targets — public pride means public exposure and an identifiable owner in the comments.
Vendor case-study pages	Spanish chatbot and CX vendors publish customer logos. Every logo is a qualified target, and the vendor may later become a partner.
The companies' own websites	The decisive check. Open the site, find the chat widget, confirm it answers substantive questions rather than routing to a form.
LinkedIn people search	The owner. Titles: "Responsable de Atención al Cliente", "Director de Transformación Digital", "Responsable de Canal Digital", "Director de Postventa", "Responsable de Servicio Técnico".

The list schema

```
company · sector · employees · city · assistant_name · assistant_url
· assistant_type (web widget | whatsapp | app | internal)
· launch_evidence_url · launch_date_approx
· owner_name · owner_title · owner_linkedin · email_pattern_guess
· tos_checked (y/n) · automation_allowed (y/n/unclear)
· probe_status (todo | running | done | skipped)
· failures_found (n) · teaser_sent_at · reply_at · report_sent_at · call_at
```

Put this in the leads CRM you already built rather than a spreadsheet — it is the one place where your existing platform earns its keep on day one.

Qualification checklist — all five must be true

- 1 The assistant is reachable without an account and answers substantive questions.
- 2 The company publishes documentation detailed enough to serve as an answer key (conditions, manuals, tariffs, FAQs).
- 3 Between 50 and 500 employees. Below 50, no budget or owner; above 500, procurement.
- 4 You can name the likely owner, with a title.

- 5 Terms of use checked; if automated access is prohibited or unclear, mark it as manual-only and probe by hand.

DISQUALIFY IMMEDIATELY

Assistants that only route to a contact form (nothing to measure), companies whose documentation is not public (no answer key, no free report possible), anything requiring an account, and any company where you cannot find a named human owner. A weak target consumes the same hours as a strong one and produces nothing.

Sources

Every factual claim in this document, referenced. Verify anything you intend to repeat to a customer — regulatory dates in particular have already moved once.

EU AI Act and the Digital Omnibus

Gibson Dunn — [EU AI Act Omnibus Agreement: postponed high-risk deadlines and other key changes](#)

Cloud Security Alliance — [EU AI Act high-risk deadline pushed to December 2027](#)

DLA Piper — [The Digital AI Omnibus: deferral of high-risk AI obligations](#)

Praxikon — [Postponement of high-risk obligations to December 2027: what changes and what still applies](#)

AI Governance Desk — [Article 10 decoded: EU AI Act data governance rules for high-risk AI](#)

Spanish AI governance law and AESIA

La Moncloa — [Consejo de Ministros, 26 May 2026: regulation of AI for trustworthy, ethical use](#)

Congreso de los Diputados — [Proyecto de Ley Orgánica 121/000096 \(BOCG\)](#)

Economist & Jurist — [Ley Orgánica de IA: AESIA, sanciones y sandboxes](#)

Spanish AI adoption and public funding

IONOS / YouGov — [35% of Spanish SMEs will invest in AI in 2026](#)

Addmira — [Adopción de IA en pymes españolas: reconciling INE and EC figures](#)

Red.es — [Kit Consulting programme · adhesion requirements for digital advisers](#)

España Digital 2026 — [Programa Kit Consulting](#)

Evaluation, RL environments and the reliability gap

Epoch AI — [An FAQ on reinforcement learning environments](#)

Troveo — [RL environment companies in 2026: the landscape](#)

MarkTechPost — [Top LLM observability and evaluation platforms in 2026](#)

Latitude — [Best AI agent evaluation platforms in 2026](#)

Inovabeing — [Why AI agents fail in production: the reliability gap in 2026](#)

Dellon Stefanus — [Most enterprise AI agents never make it to production in 2026](#)

Golden datasets and evaluation practice

Microsoft Data Science — [The path to a golden dataset, or how to evaluate your RAG](#)

Statsig — [Golden datasets: creating evaluation standards](#)

Maxim AI — [Building a golden dataset for AI evaluation](#)

QASkills — [RAG regression testing: preventing quality drift](#)

Survey figures on pilot-to-production rates vary widely by methodology and by how "production" is defined; they are used here as directional evidence of a reliability gap, not as precise measurements. Regulatory dates were accurate as of early September 2026 — re-verify before citing them to a customer.