

MANAGED DATA PROCUREMENT FOR AI TEAMS

# From landing page to first paid dataset clients.

A practical operating guide for selling Caudals as a B2B data sourcing, licensing, preprocessing, and delivery service before scaling the platform.

14

days to buyer signal

The first milestone is not more software. It is 15 buyer discovery calls.

€2.5k

first paid sprint floor

Start with a feasibility sprint before committing to a full dataset build.

1

narrow wedge

Spanish and EU business datasets for AI support, ecommerce, search, and evaluation.

Buyer promise

Tell us the dataset your model needs. We find sources, negotiate access, clean it, verify quality, and deliver it.

Supplier promise

Turn unused business data into controlled licensed revenue, with preprocessing and risk review handled.

Build principle

Sell manually first. Automate only the steps that repeat across paid pilots.

# Start as a managed service, not a marketplace.

Caudals already has the platform foundations for requester briefs, supplier workflows, admin moderation, exports, payments, and support intake. The startup risk is not whether you can build screens. The risk is whether buyers will pay for a specific data outcome.

**First milestone:**  
Book 15 buyer discovery calls, sell one paid feasibility sprint, and use that pilot to decide the next product automation.

POSITIONING

What Caudals sells

Caudals helps AI companies source, license, clean, document, and deliver high-quality business datasets from companies and public sources.

This is easier to buy than a generic dataset marketplace because the buyer has one acute problem: a model, agent, search system, or evaluation process is blocked by missing data.

AVOID EARLY

Do not begin in the hardest categories

Avoid healthcare, children, biometrics, precise location, financial eligibility, raw personal communications, or other highly sensitive categories for the first pilots.

Begin with business datasets where rights, anonymization, and buyer value can be tested with lower operational risk.

The first version of Caudals should feel like a specialist data procurement team with a secure portal, not an empty catalog.

Buyer side

Custom dataset requests, feasibility assessment, sourcing, preprocessing, QA, documentation, and delivery.

Supplier side

Data assessment, anonymization review, licensing structure, approved-use controls, and revenue share.

Admin side

Lead pipeline, deal tracking, rights checklist, source registry, quality reports, export controls, and audit logs.

# Win a narrow use case before broadening the platform.

The best first wedge is Spanish and EU business data for AI support, ecommerce search, product enrichment, and model evaluation. It fits Caudals' EU context, has visible buyer pain, and can use a mix of company data plus public sources.

**Rule:**  
A good first dataset has a buyer who can describe exactly how they will test it.

IDEAL BUYER PROFILE

### Who to target first

AI startupsVertical SaaSCommerce AI

Support automationIndustrial AILegal-techRetail-tech

Logistics AI

Look for teams with a live AI feature, a model release deadline, or a known data quality problem.

BUYER TITLES

### Who owns the problem

- Founder or CEO at a seed-stage AI company.
- CTO or Head of Engineering adding AI features.
- Head of AI, Applied ML Lead, or Data Lead.
- Product Lead for AI agents, search, support, or analytics.
- Data Partnerships or Data Acquisition lead, if the company is larger.

## Buying signals to search for

| SIGNAL                                                    | WHAT IT MEANS                                                | SEARCH PATTERN                                      |
|-----------------------------------------------------------|--------------------------------------------------------------|-----------------------------------------------------|
| Hiring for fine-tuning, RAG, LLM eval, or data annotation | They have an active model program and a data bottleneck.     | "LLM evaluation" "data" "Spain"                     |
| Launching Spanish or EU AI features                       | They need local language, domain, and compliance-aware data. | "AI support" "Spanish" "startup"                    |
| Recent funding                                            | They have budget and urgency to ship product milestones.     | Crunchbase, Sifted, Tech.eu, LinkedIn funding posts |
| Public demo with weak domain coverage                     | They may need edge cases, evaluation sets, or expert data.   | Product Hunt, demo videos, launch posts             |

# Sell three concrete products.

Each offer should have a specific output, timeline, and acceptance point. This lets buyers say yes without requiring a long enterprise procurement process.

**Pricing logic:**  
Charge for expertise, sourcing difficulty, legal review, cleaning, QA, and delivery. Do not price like raw rows in a spreadsheet.

|                                                                                                                                                                                                                                                                                                      |                                                                                                                                                                                                                                                                           |                                                                                                                                                                                                                                                                    |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <div>OFFER 01</div> <div><b>Dataset Feasibility Sprint</b></div> <div><b>Price:</b> €2,500 to €5,000</div> <div><b>Timeline:</b> 5 to 7 business days</div> <div>Deliver a dataset brief, source map, supplier shortlist, compliance risk map, sample schema, quality rubric, and pilot quote.</div> | <div>OFFER 02</div> <div><b>Pilot Dataset Build</b></div> <div><b>Price:</b> €12,000 to €35,000</div> <div><b>Timeline:</b> 3 to 5 weeks</div> <div>Deliver a real sample pack, cleaned dataset, provenance table, QA report, dataset card, and buyer-ready export.</div> | <div>OFFER 03</div> <div><b>Private Data Supply</b></div> <div><b>Price:</b> €3,000 to €15,000/month</div> <div><b>Timeline:</b> recurring</div> <div>Maintain refreshed datasets, evaluation sets, source relationships, quality checks, and delivery SLAs.</div> |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

## Acceptance criteria for a pilot

| DIMENSION | EXAMPLE BUYER ACCEPTANCE RULE                                                         | CAUDALS PROOF                                                  |
|-----------|---------------------------------------------------------------------------------------|----------------------------------------------------------------|
| Rights    | Data can be used for AI training or evaluation under the agreed scope.                | Signed supplier license, source inventory, approved-use terms. |
| Quality   | Less than 2 percent duplicate records, required fields above 95 percent completeness. | Automated quality report and sample review notes.              |
| Coverage  | Includes at least 20 categories, Spanish text, and meaningful edge cases.             | Distribution analysis, category counts, manual spot checks.    |
| Format    | Delivered as JSONL plus Parquet with stable field names and metadata.                 | Export package, schema file, dataset card.                     |
| Safety    | No unapproved personal data or restricted fields in the final export.                 | PII scan, redaction log, human release approval.               |

# The next 14 days are sales execution.

Your goal is to force a clear market answer quickly. A finished landing page is enough to begin. Build a focused lead list, send personalized outreach, and convert the strongest pain into paid feasibility sprints.

## Output target:

100 buyer leads, 20 personalized emails per day, 15 discovery calls, 1 paid sprint.

**Day 1** Choose the first wedge, write one buyer page, one supplier page, and one sample dataset card.

**Days 2-3** Build a list of 100 buyers from LinkedIn, job posts, launch directories, AI company lists, funding news, and Hugging Face organizations.

**Days 4-10** Send 20 personalized emails per day. Track reply rate, meeting rate, pain category, budget signal, and urgency.

**Days 7-14** Run discovery calls. Ask for paid sprint commitment when the data problem is specific and urgent.

## Buyer email script

Subject: Dataset for [specific AI use case]

Hi [Name],

I saw [specific signal: your AI support launch / hiring for ML data / Spanish expansion].

I am building Caudals, a managed data procurement service for AI teams. We help companies source, license, clean, and document domain-specific datasets from business sources and public data.

For [Company], I suspect [specific dataset] could improve [model, eval, or product outcome].

Would it be useful to compare what data you need for the next model release and whether we can source a pilot dataset in 3 to 4 weeks?

## Discovery call questions

- What AI feature or model is blocked by data?
- What data do you already have?
- What is missing: volume, quality, diversity, labels, freshness, language, geography, edge cases?
- What exact test would make the dataset accepted?
- What data must not be used?
- If the sourcing plan is credible, can you pay for a pilot this month?

# Turn unused company data into controlled licensed inventory.

Do not ask suppliers to join a vague marketplace. Ask for a low-friction assessment of whether their operational data can become a controlled, licensed dataset.

**Trust comes first:**  
Early suppliers should approve every buyer and every licensed use case until the process is proven.

SUPPLIER TARGETS

Where useful data lives

- Ecommerce companies with product catalogs, FAQs, returns reasons, or category taxonomies.
- BPO and support operations with anonymizable ticket taxonomies.
- Industrial firms with manuals, maintenance logs, defects, or parts data.
- Vertical SaaS companies with aggregated workflow data.
- Marketplaces, associations, agencies, and data-rich SMEs.

DEAL STRUCTURES

How to make it easy to say yes

|         |                                                   |
|---------|---------------------------------------------------|
| SIMPLE  | One-time license fee                              |
| ALIGNED | Revenue share, often 30 to 60 percent to supplier |
| SAFER   | Non-exclusive licensing first                     |
| TRUSTED | Buyer-by-buyer approval for first deals           |

Supplier email script

Subject: Revenue from unused [data type]

Hi [Name],

Many companies have operational datasets that are valuable for AI training and evaluation but are not packaged for licensing.

Caudals helps companies assess, clean, anonymize when needed, and license approved datasets to AI companies under controlled terms. You keep approval over what is shared and with whom.

Would you be open to a 20-minute review of whether your [catalog / support / industrial / logistics] data could become a licensed dataset?

Supplier qualification checklist

- Who owns the data and who can grant rights?
- Was the data collected under terms that allow licensing or AI use?
- Does the data include personal, confidential, copyrighted, or restricted fields?
- Can a sample be reviewed under NDA?
- Does the supplier want one-time revenue, recurring revenue, or buyer-by-buyer approval?

# Deliver the first pilots with a rigorous manual workflow.

Before investing in complex platform features, create a repeatable operating system for one dataset request. The software should mirror this workflow later.

## Human approval stays mandatory:

Rights review, sensitive-data release, final buyer delivery, and commercial terms should not be fully automated.

**01 Intake** Capture buyer use case, model need, data type, geography, language, timeline, budget, sensitivity, and acceptance criteria.

**02 Brief** Translate the call into a dataset brief with target schema, source hypotheses, exclusions, quality rubric, and delivery format.

**03 Source** Map suppliers and public sources. Contact suppliers with a specific ask tied to a real buyer use case.

**04 Contract** Sign NDA, data evaluation agreement, supplier license, buyer license, and DPA where required.

**05 Ingest** Receive data through secure upload, hash files, record source metadata, and isolate raw data from cleaned exports.

**06 Process** Normalize, deduplicate, validate schema, scan for sensitive data, redact where approved, and create quality reports.

**07 Deliver** Package the dataset, dataset card, schema, source inventory, QA report, and limitations statement.

## Minimum delivery package

- Cleaned dataset in JSONL, Parquet, and CSV if needed.
- Schema and field dictionary.
- Dataset card with intended use, limitations, provenance, and known risks.
- Quality report with completeness, duplicates, distributions, and manual review notes.
- License summary and approved-use restrictions.

## Internal evidence to keep

- Supplier permission and contract references.
- Raw file hashes and processing logs.
- Sensitive-data scan results and release approval.
- Buyer acceptance criteria and sign-off.
- Incident and takedown process.

# Build the smallest portal around the managed service.

The existing Caudals stack is enough for the next stage. Add two public intake paths, a private admin pipeline, and processing automation around actual paid work.

**Current stack fit:**  
Next.js, Supabase, Stripe, Resend, and DigitalOcean Spaces already match the service workflow.

## What to build now versus later

| NOW                                  | WHY NOW                                                                    | LATER                                                         |
|--------------------------------------|----------------------------------------------------------------------------|---------------------------------------------------------------|
| Request a dataset intake form        | Converts landing traffic and outreach replies into qualified buyer briefs. | Self-serve request wizard with dynamic pricing.               |
| Offer data intake form               | Starts supplier pipeline and source registry.                              | Supplier portal with contract workflow and revenue dashboard. |
| Admin lead/deal pipeline             | Lets you manage buyer and supplier conversations without spreadsheets.     | Automated CRM sync and scoring model.                         |
| Secure upload and controlled exports | Required for paid pilot delivery.                                          | Client-specific data rooms and API access.                    |
| Quality report generator             | Creates enterprise trust and repeatable delivery evidence.                 | Buyer-facing live data observability dashboard.               |

|                                                                                                                                                                                                                                                              |                                                                                                                                                                                                                                                               |                                                                                                                                                                                                                                                                                           |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <div>PROCESSING</div> <div>Recommended tools</div> <ul style="list-style-type: none"><li>• Python for pipelines.</li><li>• Polars and DuckDB for structured processing.</li><li>• Parquet for canonical storage.</li><li>• JSONL for LLM datasets.</li></ul> | <div>QUALITY</div> <div>Recommended tools</div> <ul style="list-style-type: none"><li>• Pandera or Great Expectations.</li><li>• Custom acceptance checks per buyer.</li><li>• Manual sample review for release.</li><li>• Dataset card generation.</li></ul> | <div>GOVERNANCE</div> <div>Recommended tools</div> <ul style="list-style-type: none"><li>• Microsoft Presidio for local PII detection.</li><li>• Google Sensitive Data Protection or AWS Macie later.</li><li>• DVC or lakeFS for versioning.</li><li>• Audit logs in Supabase.</li></ul> |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

## Automation backlog

| AUTOMATION          | TRIGGER                        | OUTPUT                                                          |
|---------------------|--------------------------------|-----------------------------------------------------------------|
| Lead scoring        | New buyer or supplier intake   | Urgency, budget, data sensitivity, and fit score.               |
| Brief generator     | Discovery notes saved          | Dataset brief, acceptance criteria, and source hypotheses.      |
| Schema profiler     | Raw data uploaded              | Field types, nulls, duplicates, distributions, and row counts.  |
| Sensitive-data scan | Raw or cleaned dataset updated | PII/sensitive field report and release blockers.                |
| Export packager     | Admin approves release         | ZIP with dataset, schema, card, license summary, and QA report. |

# Make trust a product feature from day one.

Dataset procurement becomes risky when rights, personal data, intended use, or sublicensing are vague. Caudals should treat provenance and permitted use as first-class product data.

**Not legal advice:**  
Get counsel before monetizing personal, sensitive, regulated, or high-risk datasets.

TRACK PER DATASET

Required dataset ledger fields

- Source owner and commercial contact.
- Permission to license, AI-training rights, and sublicensing restrictions.
- Collection context and original terms.
- Personal data, confidential data, IP, and sensitive category flags.
- Approved uses, prohibited uses, retention rules, and deletion obligations.
- Buyer-specific access history and export versions.

CONTRACTS

Minimum legal pack

- Mutual NDA for early samples.
- Data evaluation agreement for pre-purchase review.
- Supplier data license with rights warranties and approved uses.
- Buyer data license with scope, restrictions, and no resale unless agreed.
- DPA when personal data is involved.
- Security exhibit for enterprise buyers.

Practical guardrails

| RISK                     | DEFAULT CAUDALS RULE                                                                      | EVIDENCE                                     |
|--------------------------|-------------------------------------------------------------------------------------------|----------------------------------------------|
| Personal data            | Prefer non-personal or truly anonymized business data for first pilots.                   | PII scan, human review, contract scope.      |
| Re-identification        | Treat pseudonymized data as still risky unless counsel approves the use case.             | Redaction method, residual risk note.        |
| Unauthorized AI training | Do not accept supplier data unless AI use and sublicensing are explicitly covered.        | Supplier license and approved-use statement. |
| Buyer misuse             | Ban re-identification, unauthorized resale, and use outside the licensed purpose.         | Buyer license and audit logs.                |
| Regulatory drift         | Review EU AI Act, GDPR, Data Governance Act, Data Act, and local rules for each category. | Risk memo attached to dataset record.        |

# Three examples you can sell.

Use these examples to make outreach concrete. Each one has a buyer pain, supplier path, deliverable, and pricing range.

**Sales test:**  
If a prospect cannot name how the dataset improves a model, defer the pilot.

|                                                                                                                                                                                                                                                                                                                                                                                                                                          |                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <p><b>EXAMPLE 01</b></p> <p><b>Ecommerce product enrichment</b></p> <p><b>Buyer:</b> AI shopping/search startup.</p> <p><b>Need:</b> 50k Spanish product records with attributes, descriptions, taxonomy, and edge cases.</p> <p><b>Sources:</b> Retailers, PIM vendors, public catalogs.</p> <p><b>Deliver:</b> Normalized product data, taxonomy, duplicate detection, QA report.</p> <p><b>Price:</b> €5k sprint plus €18k pilot.</p> | <p><b>EXAMPLE 02</b></p> <p><b>Support-ticket intent dataset</b></p> <p><b>Buyer:</b> SaaS company training Spanish support agents.</p> <p><b>Need:</b> 20k anonymized tickets labeled by intent, urgency, and resolution type.</p> <p><b>Sources:</b> BPOs, ecommerce support teams, SaaS support operations.</p> <p><b>Deliver:</b> JSONL conversations, PII scan, label guide, eval split.</p> <p><b>Price:</b> €25k to €40k, with higher compliance burden.</p> | <p><b>EXAMPLE 03</b></p> <p><b>Industrial maintenance corpus</b></p> <p><b>Buyer:</b> Predictive maintenance AI company.</p> <p><b>Need:</b> Manuals, fault reports, maintenance logs, part descriptions.</p> <p><b>Sources:</b> Maintenance contractors, equipment distributors, public manuals.</p> <p><b>Deliver:</b> Clean document corpus, metadata, OCR where needed, retrieval/eval set.</p> <p><b>Price:</b> €20k to €60k depending on rights and processing.</p> |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

## What to say on the call

| IF THEY SAY...                       | RESPOND WITH...                                                                                                                               |
|--------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------|
| "We are not sure what data we need." | "That is exactly the feasibility sprint. We turn your product goal into a concrete dataset spec and sourcing plan."                           |
| "Can you guarantee the data exists?" | "Not before sourcing. We can guarantee the feasibility process, supplier outreach, risk map, and a go/no-go recommendation."                  |
| "We need this fast."                 | "Then we should narrow scope to a pilot dataset with a clear acceptance test, not try to source the full production corpus first."            |
| "We need personal data."             | "We can assess that, but the first step is legal basis, rights, minimization, anonymization, and buyer restrictions before any sample moves." |

# Move from calls to repeatable revenue.

The plan is intentionally commercial first. Product work happens in parallel, but only around the pipeline and delivery steps that make paid pilots easier to close.

## North-star metric:

Paid dataset pipeline value, not waitlist count.

**Week 1** Pick wedge, update landing CTAs, create buyer/supplier one-pagers, create sample dataset card, build 200-lead list.

**Weeks 2-3** Run outbound, book discovery calls, qualify pain, ask for paid feasibility sprint, begin supplier outreach for repeated needs.

**Weeks 4-5** Sell 1 to 2 feasibility sprints. Convert one into a pilot dataset build. Deliver manually with strong documentation.

**Days 45-60** Automate intake scoring, brief generation, schema profiling, quality reports, and export packaging around the paid pilot.

**Days 60-90** Publish two anonymized case studies, build supplier registry, create reusable templates, and raise prices after proof.

## This week checklist

- ☐ Choose the first wedge and write it on the landing page.
- ☐ Replace generic waitlist CTA with "Request a dataset" and "Offer data".
- ☐ Add budget, timeline, data type, geography, and sensitivity fields.
- ☐ Create a simple buyer lead sheet and supplier lead sheet.
- ☐ Send the first 40 personalized buyer emails.

## Do not spend this month on

- ☐ A public dataset marketplace with no supply.
- ☐ Complex self-serve pricing before you know buyer willingness to pay.
- ☐ A full CRM when a simple pipeline is enough.
- ☐ High-risk regulated datasets before legal process is mature.
- ☐ Contributor-scale features before supplier and buyer demand is validated.

## Reference sources

1. [European Commission: GDPR rules for businesses and organisations](#)
2. [European Commission: Data Governance Act](#)
3. [European Commission: Data Act](#)
4. [EUR-Lex: Regulation \(EU\) 2024/1689, AI Act](#)
5. [Microsoft Presidio documentation](#)
6. [Google Cloud Sensitive Data Protection: de-identification](#)
7. [Great Expectations documentation](#)
8. [DuckDB Parquet documentation](#)
9. [DVC documentation](#)
10. [lakeFS documentation](#)