

Training data for AI, *taylor-made*.

We connect companies that hold large volumes of valuable data with AI teams that need real, sector-specific information from the industries they operate in.

What is Caudals?

We connect those who hold the data with those who *need it*.

Caudals is a data marketplace built for companies developing artificial intelligence models that need real-world information from the industries they work in. We connect companies holding valuable data with teams that need it to train their models, and we make sure the transfer is structured, with clear permissions and in a format AI can actually read.

The problem we solve

An analogy that makes it easy.

AI is the chef. Data is the ingredients.

Without good ingredients, not even the best chef puts out a great dish. Today almost everyone is competing to have better chefs, while almost no one is taking care of the supplies. That is where Caudals comes in: we take care of sourcing the raw material AI needs in order to learn.

Good ingredients (real-world, sector-specific data) are rarely found on the open internet or in public databases. They live inside companies that generate them as a byproduct of their work, but that are not in the business of selling them. The data exists, but it is scattered and hard to obtain.

AI teams spend months negotiating with each individual owner, or worse, never get to train the model at all. Caudals exists so that this step stops being a bottleneck.

How we work

Three steps.

01 We identify

We find which companies hold the data a specific AI project requires, and we approach them.

02 We agree

We sit down with the owner and define what is shared, under what conditions, for what use, and at what price. The owner always has the final say.

03 We deliver

We prepare the data in the exact format the AI team will use (labeled images, structured tables, audio with transcripts, etc.) and hand it over ready to work with.

Why it matters if your company holds data

A simple way to capture value from what you already have.

Many companies accumulate, over years, data they generate as part of their normal operations: images, measurements, historical records, logs, archives, and so on. Data that holds real value for third parties, but that is never monetized because it is not the company's core business, and there is no time to build the commercial channel needed to turn it into revenue.

You decide what is shared

There is no need to open up your entire archive. We can begin with a very specific subset and see how it performs.

You remain the owner

Caudals receives a usage license to deliver the data to specific clients, never ownership of the information itself.

You set the terms

What the data can be used for, in what format, with what level of detail, for how long. We define it together before any transfer takes place.

You receive compensation

The model (per-volume payment, per-project fee, recurring revenue) is agreed based on whatever makes the most sense for you.

We provide the channel

Buyer sourcing, negotiation, technical preparation of the data and ongoing operational follow-up. You do not need to build anything new inside your company.

An important guarantee

We only work with data shared voluntarily.

Caudals does not scrape data from the internet on its own, nor use any information without the owner's permission. Every dataset arrives through an explicit agreement with the company that owns it, and that agreement also defines what the buyer is allowed to do: train and validate AI models, without republishing or showing the content to end users.